

# Tutorial 02 – C++ Threaded Programming

Parallel Programming 2024  
Tutorial 02



Dennis-Florian Herr, M.Sc.

Chair for Computer Architecture and Parallel Systems (CAPS)

Technical University Munich

May 07, 2024



*TUM Uhrenturm*

## Assignment 1: VV-AES

# General Ideas to optimize sequential code

- Improve algorithmic performance (think about asymptotic complexity / Big O)
- Think about the right data structure for the problem
- Optimize cache usage (cache access is orders of magnitude faster than main memory access)
- Reuse previously computed values (cf. dynamic programming)
- Consider precomputing an often recomputed small amount of data
- Do not reinvent the wheel (check *\*at least\** the standard library for existing solutions)

# Improve substitute\_bytes

Slow approach by repeatedly searching through the originalCharacter array

## Before

```
1 void substitute_bytes() {
2     // For each byte in the message
3     for (int column = 0; column < BLOCK_SIZE; column++) {
4         for (int row = 0; row < BLOCK_SIZE; row++) {
5             int index = -1;
6             for (int i = 0; i < UNIQUE_CHARACTERS; i++) {
7                 if (originalCharacter[i] == message[row][column]) {
8                     index = i;
9                 }
10            }
11            message[row][column] = substitutedCharacter[index];
12        }
13    }
14 }
```


Create a lookup map reducing the search from up to 256 iterations to 1

## Optimized

```
1  uint8_t substituteMap[UNIQUE_CHARACTERS];
2  void create_substitute_map(){
3      for(int i = 0; i < UNIQUE_CHARACTERS; i++){
4          substituteMap[originalCharacter[i]] = substitutedCharacter[i];
5      }
6  }
7  void substitute_bytes() {
8      // For each byte in the message
9      for (int column = 0; column < BLOCK_SIZE; column++) {
10         for (int row = 0; row < BLOCK_SIZE; row++) {
11             message[row][column] = substituteMap[message[row][column] ];
12         }
13     }
14 }
```

# Improve substitute\_bytes

And we're done.... (oops)

Runtime	Speedup	Status
8.11421	8.25712	 Passed

Don't reinvent the wheel. Use the std library rotate function

optimized

```
1  #include <algorithm>
2  void shift_rows() {
3      // Shift each row, where the row index corresponds to how many columns the data is shifted.
4      for (int row = 0; row < BLOCK_SIZE; ++row) {
5          std::rotate(std::begin(message[row]), std::begin(message[row]) + row, std::end(message[row]));
6      }
7  }
```

# Improve shift\_rows


Alone it's dwarfed by substitute\_bytes

59.94560

1.11768

 Too Slow

But together it's a significant boost

Runtime	Speedup	Status
3.77289	17.75829	 Passed



## Precompute all possible powers of 0 to 255

### optimized

```
1  int powers[256][BLOCK_SIZE + 1];
2  // precomputed powers of all possible message values (256)
3  for (int i = 0; i < 256; i++) {
4      for (int j = 1; j <= BLOCK_SIZE; j++) {
5          powers[i][j] = power(i, j);
6      }
7  }
8  void mix_columns() {
9      for (int column = 0; column < BLOCK_SIZE; ++column) {
10         for (int row = 0; row < BLOCK_SIZE; ++row) {
11             int result = 0;
12             for (int degree = 0; degree < BLOCK_SIZE; degree++) {
13                 result += polynomialCoefficients[row][degree] * powers[message[degree][column]][degree + 1];
14             }
15             message[row][column] = result;
16         }
17     }
18 }
```

# Improve substitute\_bytes

Runtime	Speedup	Status
3.04771	21.98369	✓ Passed

## Use CISC PSHUFB: Packed Shuffle Bytes

```
1  uint8_t shift_row_mask[64] = {
2      0, 1, 2, 3, 4, 5, 6, 7, 9, 10, 11, 12, 13, 14, 15, 8,
3      2, 3, 4, 5, 6, 7, 0, 1, 11, 12, 13, 14, 15, 8, 9, 10,
4      4, 5, 6, 7, 0, 1, 2, 3, 13, 14, 15, 8, 9, 10, 11, 12,
5      6, 7, 0, 1, 2, 3, 4, 5, 15, 8, 9, 10, 11, 12, 13, 14
6  };
7
8  inline void shift_rows() {
9      __asm__ (
10         "leaq    shift_row_mask(%rip), %rax\n"
11         "vmovdqa message(%rip), %xmm0\n"
12         "vpshufb (%rax), %xmm0, %xmm0\n"
13         "vmovaps %xmm0, message(%rip)\n"
14         "vmovdqa 16+message(%rip), %xmm1\n"
15         "vpshufb 16(%rax), %xmm1, %xmm1\n"
16         "vmovaps %xmm1, 16+message(%rip)\n"
17         "vmovdqa 32+message(%rip), %xmm2\n"
18         "vpshufb 32(%rax), %xmm2, %xmm2\n"
19         "vmovaps %xmm2, 32+message(%rip)\n"
20         "vmovdqa 48+message(%rip), %xmm3\n"
21         "vpshufb 48(%rax), %xmm3, %xmm3\n"
22         "vmovaps %xmm3, 48+message(%rip)\n"
23     );
24 }
```

## A word on the leaderboard

- Friendly competition. Completely optional!
- Code/algorithm can be changed freely (as long as the result is accepted)
- ...but please don't try to hack the system!
- ...and don't write scripts to submit hundreds of times (I will delete and disqualify these)

## How to compete?

- Find optimal parallelization
- Optimize sequential performance
- Skip unnecessary work
- Approximate solutions might be good enough!

## *Need for Speed*

Share your cool solutions (after the deadline)

<https://zulip.in.tum.de/#narrow/stream/2330-ParProg-24/topic/Need.20for.20Speed>

## Theory

# C++ Threads : Passing Data

## Creating Threads

```
1      #include <thread>
2      std::thread thread_object( callable )
3      std::thread thread_object( callable , arguments)
```

# C++ Threads : Passing Data

## Creating Threads

```
1      #include <thread>
2      std::thread thread_object( callable )
3      std::thread thread_object( callable , arguments)
```

---

## Pass by Pointer

### Main

```
1  ...
2  int argument = 1;
3  std::thread thread( kernel, &argument )
4  ...
```

### Kernel

```
1  void* kernel( int* args) {
2      ++(*args);
3      ...
4  }
```



# C++ Threads : Passing Data

## Creating Threads

```
1      #include <thread>
2      std::thread thread_object( callable )
3      std::thread thread_object( callable , arguments)
```

---

## Pass by Pointer

### Main

```
1  ...
2  int argument = 1;
3  std::thread thread( kernel, &argument )
4  ...
```

### Kernel

```
1  void* kernel( int* args) {
2      ++(*args);
3      ...
4  }
```

---

## Pass by Reference

### Main

```
1  ...
2  int argument = 1;
3  std::thread thread( kernel, std::ref(argument) )
4  ...
```

### Kernel

```
1  void* kernel( int &args) {
2      ++args;
3      ...
4  }
```

## Joining Threads

```
1      #include <thread>
2      std::thread::join()
```

## Joining Threads

```
1      #include <thread>
2      std::thread::join()
```

## Example

### Main

```
1  ...
2  // Fork a thread
3  std::thread t(my_kernel);
4
5  // Join the thread
6  t.join();
7  ...
```

### Kernel

```
1  void my_kernel () {
2      sleep(1);
3      ...
4  }
```

# C++ Threads: Pitfalls

## Example 1

### Main

```
1 ...
2 // Start numThreads threads
3 std::thread threads[numThreads];
4 for (int i = 0 ; i< numThreads ; ++i){
5     threads[i] = std::thread(oops_kernel, &i);
6 }
7 for (int i = 0 ; i< numThreads ; ++i){
8     threads[i].join();
9 }
10 ...
```

### Kernel

```
1 void oops_kernel (int *args) {
2     int argument = *args;
3     ...
4 }
```

# C++ Threads: Pitfalls

## Example 1

### Main

```
1  ...
2  // Start numThreads threads
3  std::thread threads[numThreads];
4  for (int i = 0 ; i< numThreads ; ++i){
5      threads[i] = std::thread(oops_kernel, &i);
6  }
7  for (int i = 0 ; i< numThreads ; ++i){
8      threads[i].join();
9  }
10 ...
```

### Kernel

```
1 void oops_kernel (int *args) {
2     int argument = *args;
3     ...
4 }
```

✗ Use of freed memory / Race Condition

# C++ Threads: Pitfalls

## Example 1

### Main

```
1 ...  
2 // Start numThreads threads  
3 std::thread threads[numThreads];  
4 for (int i = 0 ; i< numThreads ; ++i){  
5     threads[i] = std::thread(oops_kernel, &i);  
6 }  
7 for (int i = 0 ; i< numThreads ; ++i){  
8     threads[i].join();  
9 }  
10 ...
```

### Kernel

```
1 void oops_kernel (int *args) {  
2     int argument = *args;  
3     ...  
4 }
```

✗ Use of freed memory / Race Condition

### Kernel

```
1 void wicked_kernel (int *args) {  
2     (*args) = 0;  
3     ...  
4 }
```

## Example 2

### Main

```
1  ...
2  // Start numThreads threads
3  int ids[numThreads];
4  std::thread threads[numThreads];
5  for (int i = 0 ; i< numThreads ; ++i){
6      ids[i] = i;
7      threads[i] = std::thread(a_kernel, &ids[i]);
8  }
9  for (int i = 0 ; i< numThreads ; ++i){
10     threads[i].join();
11 }
12 ...
```

### Kernel

```
1 void a_kernel (int* args) {
2     int argument = *args;
3     ...
4 }
```

## Example 2

Main

```
1  ...
2  // Start numThreads threads
3  int ids[numThreads];
4  std::thread threads[numThreads];
5  for (int i = 0 ; i< numThreads ; ++i){
6      ids[i] = i;
7      threads[i] = std::thread(a_kernel, &ids[i]);
8  }
9  for (int i = 0 ; i< numThreads ; ++i){
10     threads[i].join();
11 }
12 ...
```

Kernel

```
1 void a_kernel (int* args) {
2     int argument = *args;
3     ...
4 }
```

✓ No Problems



# C++ Threads: Pitfalls

## Example 3

### Main

```
1  std::thread threads[numThreads];
2  void lets_spawn_threads(){
3
4      int ids[numThreads];
5      for (int i = 0 ; i< numThreads ; ++i){
6          ids[i] = i;
7          threads[i] = std::thread(nooooo_kernel, &ids[i]);
8      }
9  }
10
11 int main(){
12     lets_spawn_threads();
13     for (int i = 0 ; i< numThreads ; ++i){
14         threads[i].join();
15     }
16
17 }
```

### Kernel

```
1  void noooo_kernel (int* args) {
2      int argument = *args;
3      ...
4  }
```

# C++ Threads: Pitfalls

## Example 3

### Main

```
1  std::thread threads[numThreads];
2  void lets_spawn_threads(){
3
4      int ids[numThreads];
5      for (int i = 0 ; i< numThreads ; ++i){
6          ids[i] = i;
7          threads[i] = std::thread(nooooo_kernel, &ids[i]);
8      }
9  }
10
11 int main(){
12     lets_spawn_threads();
13     for (int i = 0 ; i< numThreads ; ++i){
14         threads[i].join();
15     }
16
17 }
```


### Kernel

```
1  void noooo_kernel (int* args) {
2      int argument = *args;
3      ...
4  }
```

✗ Use of freed memory

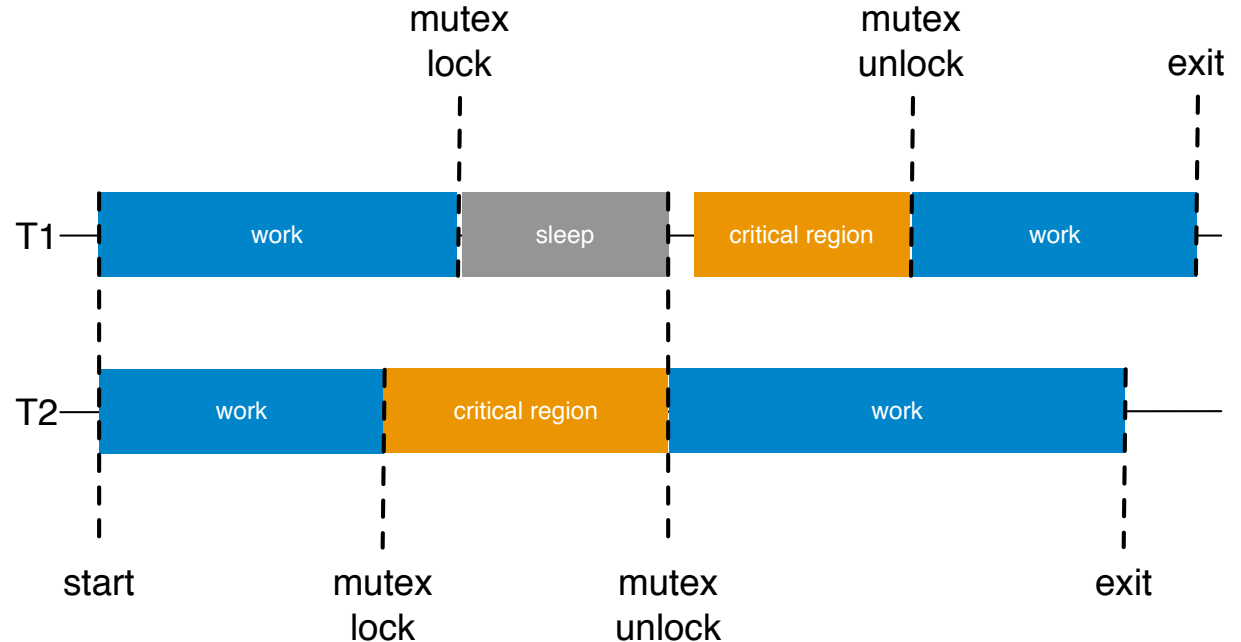
# Synchronization

- Needed for accesses to shared resources
- Drawback: Overhead
- Frequent reasons for synchronizing:
  - Prevention of concurrent access
  - Signal passing
- C++ standard library provides following mechanisms (not exhaustive):
  - Mutexes
  - Condition Variables
  - Barriers (not covered)
  - Semaphores (not covered)




# Mutex (Mutual Exclusion Lock)


- The simplest and most primitive synchronization method
- Uses atomic (hardware) operations
- Ensures absolute owner of (critical) code section
- Threads can lock and unlock mutexes



# Condition Variables – Signaling




# Condition Variables – Signaling




# Condition Variables – Signaling


**Signal**



**Broadcast**



**Timed Wait**



## Undercounting

```
1  int blockSize = SIZE / NUM_THREADS;
2
3
4
5
6  for(int i = 0; i < blockSize; i++) {
7      int dataIndex = threadId * blockSize + i;
8      ...
9  }
10 ...
```

Data



## Overcounting

```
1  int blockSize = 8;
2  int myBlock = 0;
3
4  while(myBlock * blockSize < SIZE) {
5      myBlock = getNextBlock();
6      for(int i = 0; i < blockSize; i++) {
7          int dataIndex = myBlock * blockSize + i;
8          ...
9      }
10 }
```

Data





## Moodle Quiz

<https://www.moodle.tum.de/mod/quiz/view.php?id=2976059>


## In-Class Exercise

# Help to calculate the starship's damage




The classic asteroids game

# Help to calculate the starship's damage




Our starship wants to go to the goal along the diagonal of the grid map

# Help to calculate the starship's damage




The red blocks are the intended trajectory of the starship

# Help to calculate the starship's damage




But there are many rocks floating in the space, having very complicated velocity function  
We can call `compute_vel( )` function to calculate the velocity of a rock

# Help to calculate the starship's damage



At every time step, all the rocks in the map and the starship will update their location  
We can call `update_rock( )` to update a rock's locations


# Help to calculate the starship's damage



When our starship moves towards the goal, it could crash rocks  
We can call `calc_hits()` to calculate how many rocks hit the starship at this time step




# Help to calculate the starship's damage



hits+2!

Several rocks can hits our starship at the same time

## Help to calculate the starship's damage



Help us to calculate the total hits number faster!  
Although we implemented the sequential way to calculate this, but it is too slow!  
Use C++ standard library threads to help us calculate this faster !

# Where to find the exercise ?

- Go to the following repository to get the exercise:

<https://gitlab.lrz.de/lrr-tum/teaching/parprog/ss2024/published-assignments>

- Use git to clone the exercise to your local machine:

`cd your_folder`

`git clone https://gitlab.lrz.de/lrr-tum/teaching/parprog/ss2024/published-assignments.git`

- You can pull from this repository every time a new exercise is published
- Go to folder in-class-2 for this week's task, you can also find a README.md there

# Week 2 exercise introduction:

## Use C++ standard library threads

- You will find `student_submission.cpp` the partially implemented parallel code
- Complete the `//TODO` section to
  - include necessary library
  - create threads for parallelization (pass called function, arguments etc.)
  - join threads to terminate parallelization
  - use mutex locks to avoid data racing
- achieve a speed up of 12 on the sever.
- Our server has 16 cores and 2 way hyperthreading (i.e. 32 threads)

# In-Class Exercise: Solution

# TODO#1: Include thread and mutex



```
1 // ##### TODO: include library for enabling thread and mutex #####
2 #include <thread> // For std::thread
3 #include <mutex> // For std::mutex
4 // ##### TODO END #####
```



```
1 // TODO: uncomment once you added the correct headers
2 std::thread threads[THREAD_NUM];
3 std::mutex mutex;
```

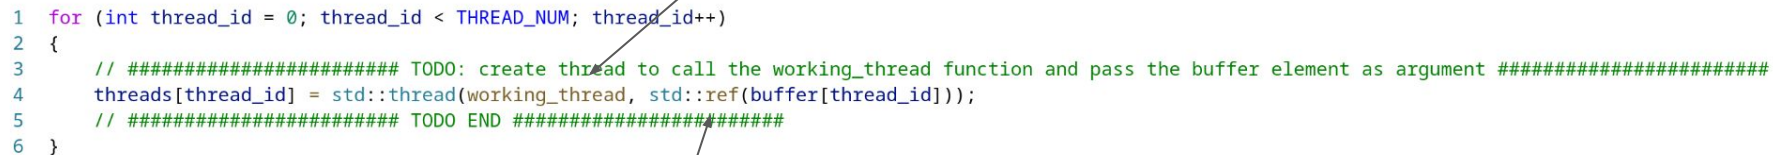
# TODO#2: Implement synchronized access to global variable



```
1 // ##### TODO: Copy task id from the global variable and increment it #####
2 // ##### DO NOT FORGET TO LOCK AND UNLOCK! #####
3 mutex.lock();
4 local_task_id = task_id++;
5 mutex.unlock();
6 // ##### TODO END #####
```

# TODO#3: Create the threads

Thread kernels need to have return type void!



```
1  for (int thread_id = 0; thread_id < THREAD_NUM; thread_id++)
2  {
3      // ##### TODO: create thread to call the working_thread function and pass the buffer element as argument #####
4      threads[thread_id] = std::thread(working_thread, std::ref(buffer[thread_id]));
5      // ##### TODO END #####
6  }
```

Ensures we pass the reference and avoids (rare) compile errors



# TODO#4: Wait for threads to complete and accumulate result




```
1  for (int thread_id = 0; thread_id < THREAD_NUM; thread_id++)
2  {
3      // ##### TODO: join thread to terminate thread, get the returned value from the buffer and add it to the (total) crashed_count #####
4      threads[thread_id].join();
5      crashed_count += buffer[thread_id];
6      // ##### TODO END #####
7  }
```

## Homework

# Assignment: Raytracer

- Task: Render a 3D scene consisting of metallic spheres using raytracing
  - The program creates a random scene based on the seed read from `stdin`.
- The renderer sends rays for each pixel.
- When a ray hits a surface it gets reflected. A ray can only be reflected a certain number of times.
- Color of the pixel is determined by the materials of the surfaces a ray hits.
- Each pixel is sampled multiple times to reduce noise.



© www.scratchapixel.com

Figure 1: Rendering with a raytracer.

- Parallelize the sequential implementation with C++ threads.
- Your speedup should be  $\geq 10$ .
- You don't have to read `maths.h` or `raytracer.h` to work on your solution.
- Evaluation command: `./student_submission -n << <seed>`
- The output is a PPM image file.
  - Try `xdg-open` or use <http://paulcuth.me.uk/netpbm-viewer> to view PPM files.




Figure 2: Raytraced scene.

# Recap & Questions

Covered today:

- Launching & Joining Threads
- Synchronization

# Questions

This slide is intentionally left blank.