

Lecture IN-2147

Parallel Programming

Martin Schulz

Technical University of Munich

Department of Computer Engineering

Lecture 2: Threading

- Seg. 6: Introduction into Parallel Architectures
- Seg. 7: Concept of Threading
- Seg. 8: The Pthreads API
- Seg. 9: Thread Synchronisation
- Seg. 10: Modern Thread APIs



Recap: Lecture 1- Introduction

Parallel processing

- Multiple tasks working together to finish a (a) larger problem (b) faster
- Goal has to be efficiency
- Multi-/Many-core developments catapult this out of a niche into every system

Programming in parallel

- Decomposition of work and data using choice of best fitting pattern
- Hybrid models are likely the best choice for individual applications
- Mapping to architectures critical

Metrics for Determining Success of Parallelization

- Speed-Up and Efficiency
- Amdahl's law to determine theoretical upper bounds
- Also here: exceptions in form of super-linear speed-ups caused by architectures

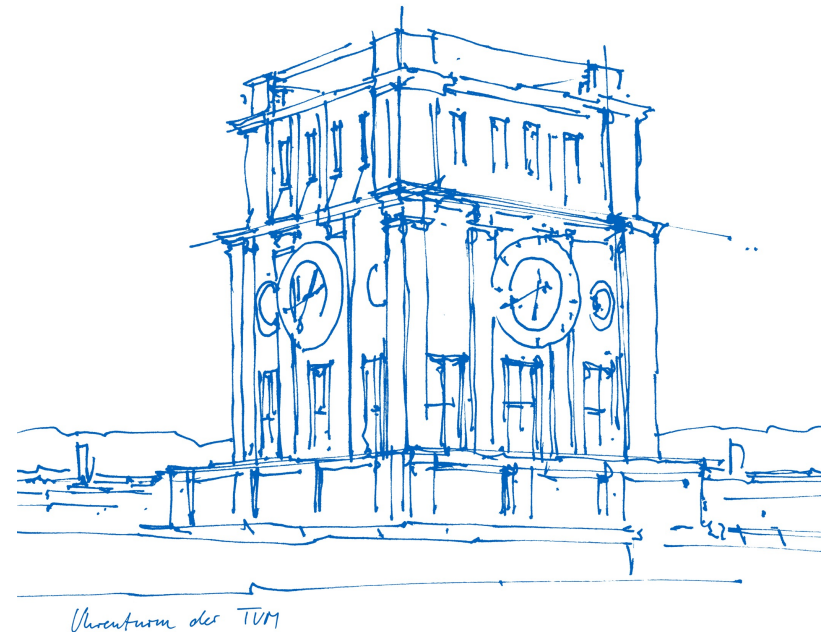
Think parallel!

Lecture IN-2147

Parallel Programming

Martin Schulz
Technical University of Munich
Department of Computer Engineering

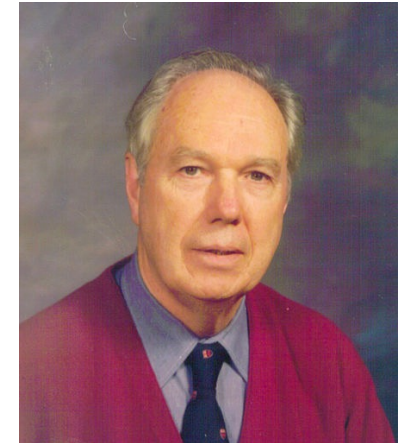
Segment 6
Introduction into Parallel Architectures



Flynn's Classification

M. Flynn, Very High-Speed Computing Systems, Proceedings of the IEEE, 54, 1966

	Single Data	Multiple Data
Single Instruction	SISD Sequential Processing	SIMD Pipelines, Vectors, GPUs
Multiple Instruction	MISD ??? / Systolic Arrays	MIMD MPP Systems Clusters




Michael Flynn
Born 1938

Relevant items covering parallel systems

- SIMD (Single Instruction Multiple Data):
Synchronized execution of the same instruction on a set of data
- MIMD (Multiple Instruction Multiple Data):
Asynchronous execution of different instructions

Classification

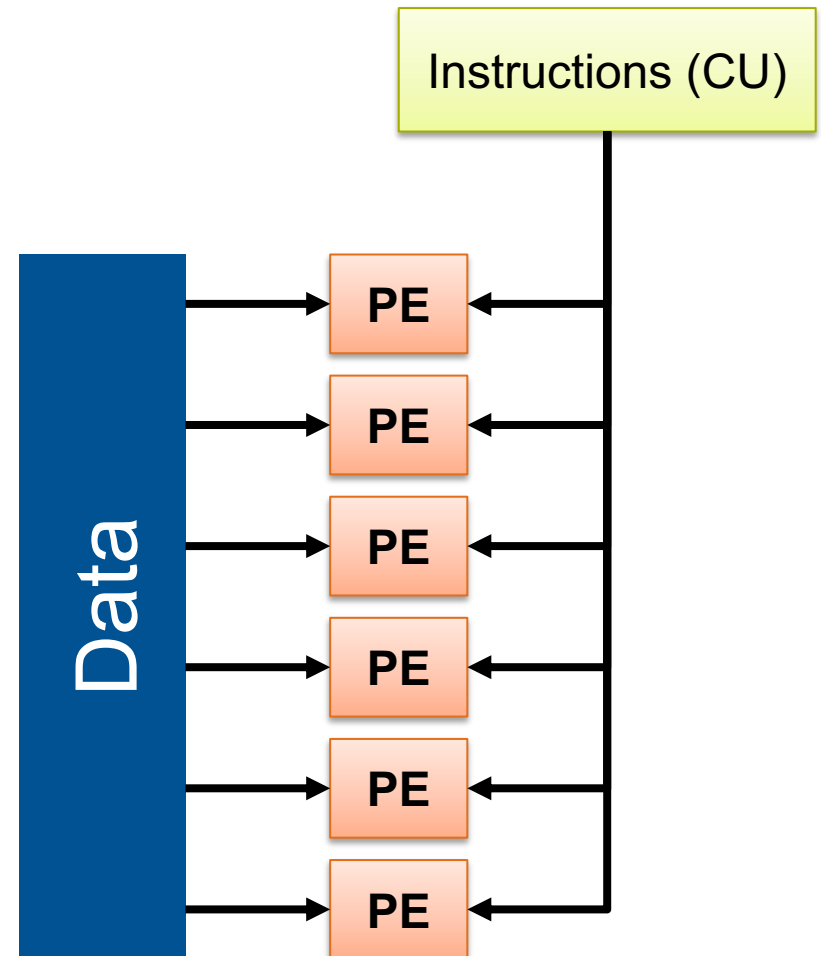


SIMD Systems

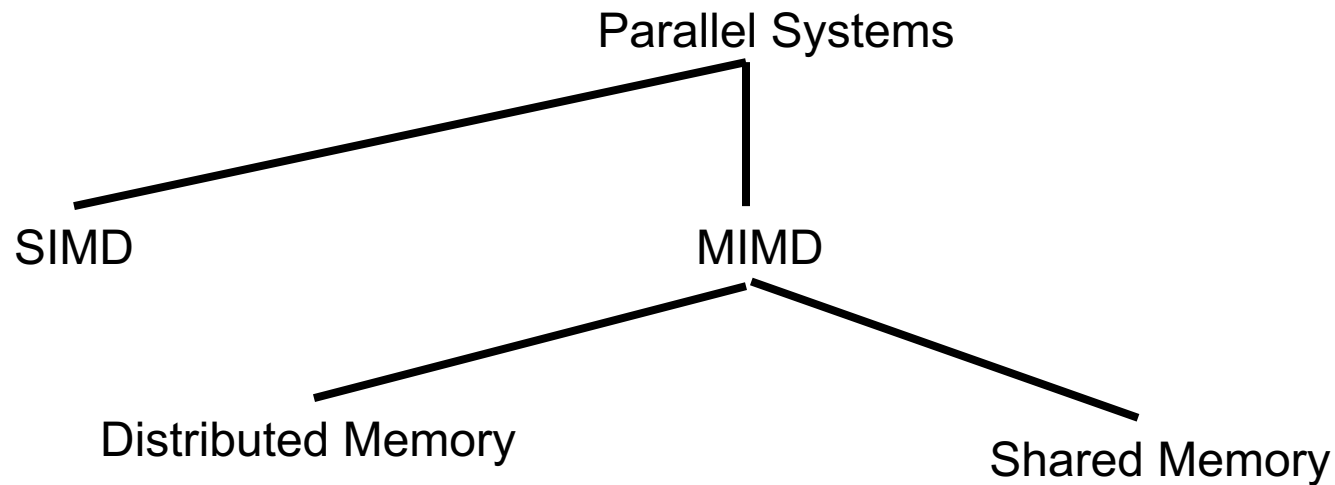
One instructions operates
on a many data streams

Vector processing
SIMD instructions

GPGPU processing
fits the same model



Classification



Shared Memory

Uniform Memory Access – UMA :
(symmetric multiprocessors - SMP):

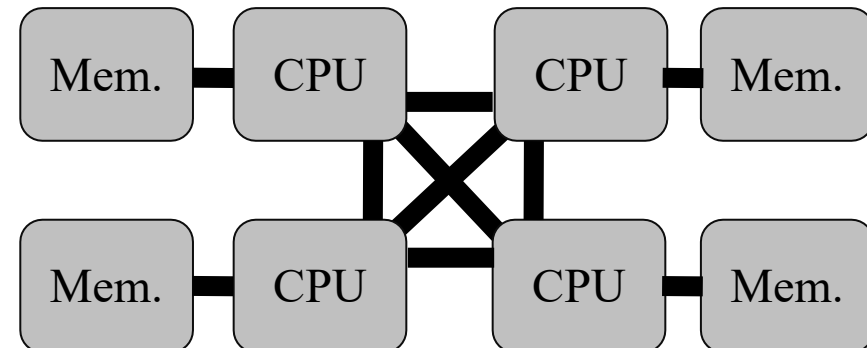
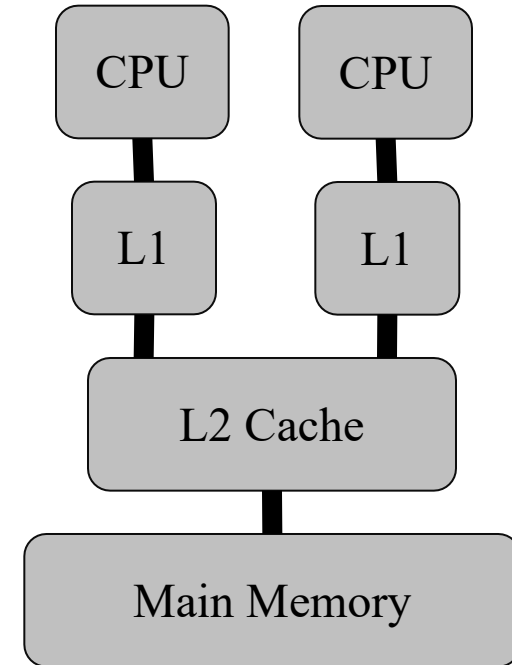
- Centralized shared memory
- Accesses to global memory from all processors have “same” latency.
- Transition from bus to crossbars

Non-uniform Memory Access Systems - NUMA
(Distributed Shared Memory Systems – HW-DSM):

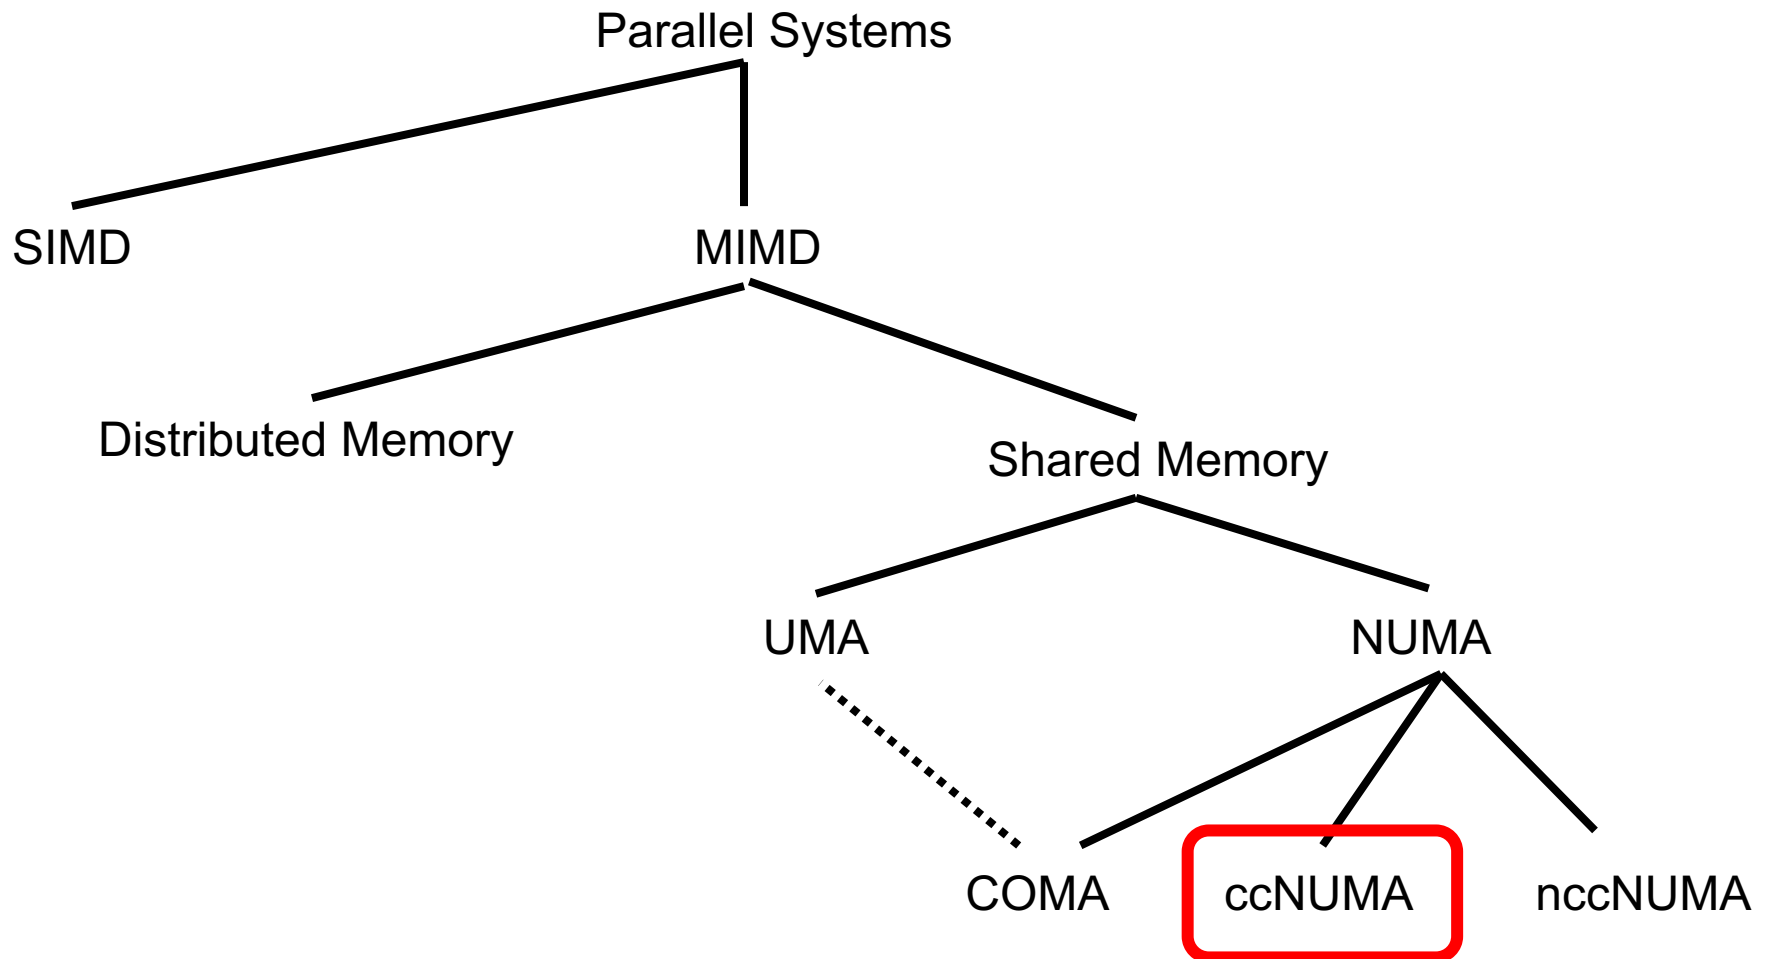
- Memory is distributed among the nodes
- Local accesses much faster than remote accesses.

More exotic

- COMA (Cacho Only)
- NCC-NUMA (non cache coherent)



Classification



Diversity in Parallel Programming Models

Driven by architecture developments (at least traditional and at the low level)

Most attached to an existing sequential programming language

- Most common: C, C++, Fortran
- Scripting languages are becoming more relevant
- APIs or language extensions

SIMD or Vector Programming

- Often in the form of pragmas (many vectorizing compilers)
- CUDA, OpenCL, ... as separate languages (but again built on base language)

Shared Memory Programming Models

- MIMD models that match shared memory architectures

Message Passing Programming Models

- MIMD models that match distributed memory architectures

Shared Memory Models Match Shared Memory

Assume a global address space with random access

- Any read/write can reach any memory cell
- This is also for NUMA systems, but locality gets tricky
- Most models assume cache coherency

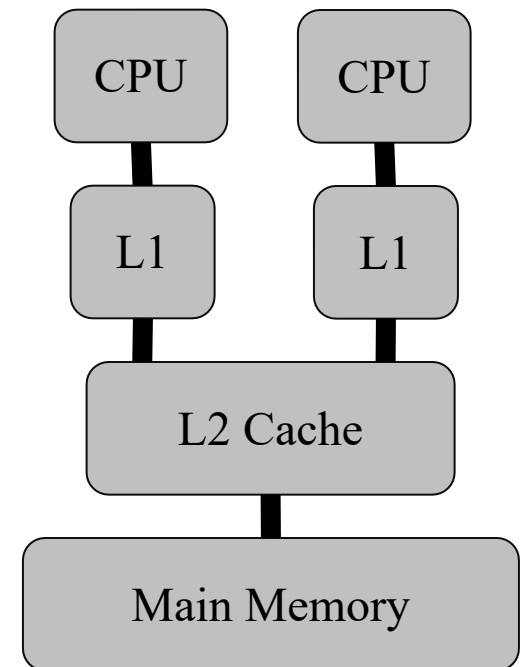
Communication through memory accesses

- Load/Store operations to arbitrary addresses
- Pass data from PE to the next

Synchronization constructs to coordinate accesses

- Need to ensure consistency
 - Data synchronization
- Need to ensure control flow
 - Control synchronization

Examples: POSIX threads, OpenMP, ...

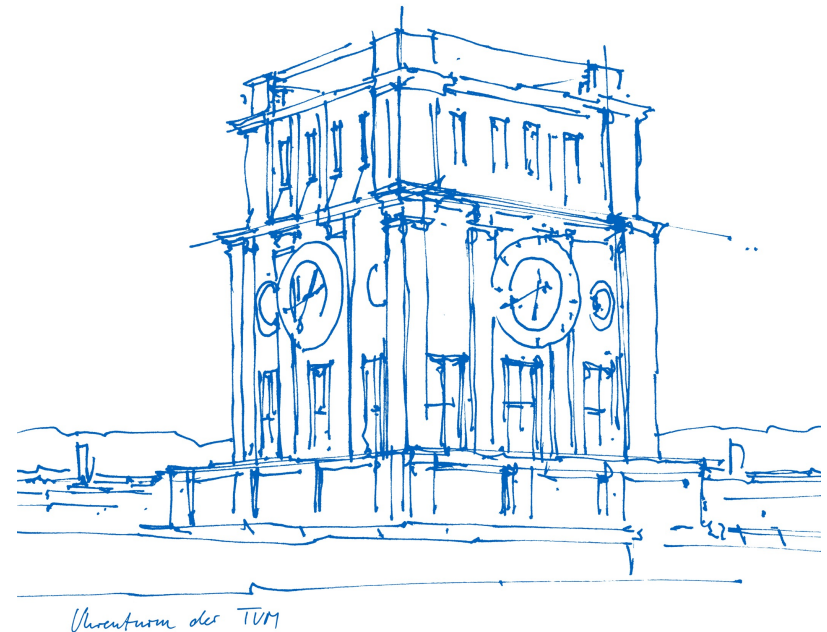


Lecture IN-2147

Parallel Programming

Martin Schulz
Technical University of Munich
Department of Computer Engineering

Segment 7
The Concept of Threading



What is a Thread?

Independent stream of execution

- Own PC
- Own Stack

Hardware threads

- Implementation of an execution stream in hardware
(think realization of a von Neumann machine in hardware)
- Separate Control Unit executing a sequence of instructions

Software threads

Hardware Threads in the Parallel Case

Traditional view

- One processor = one hardware thread (i.e., one control unit)

Boards with multiple sockets

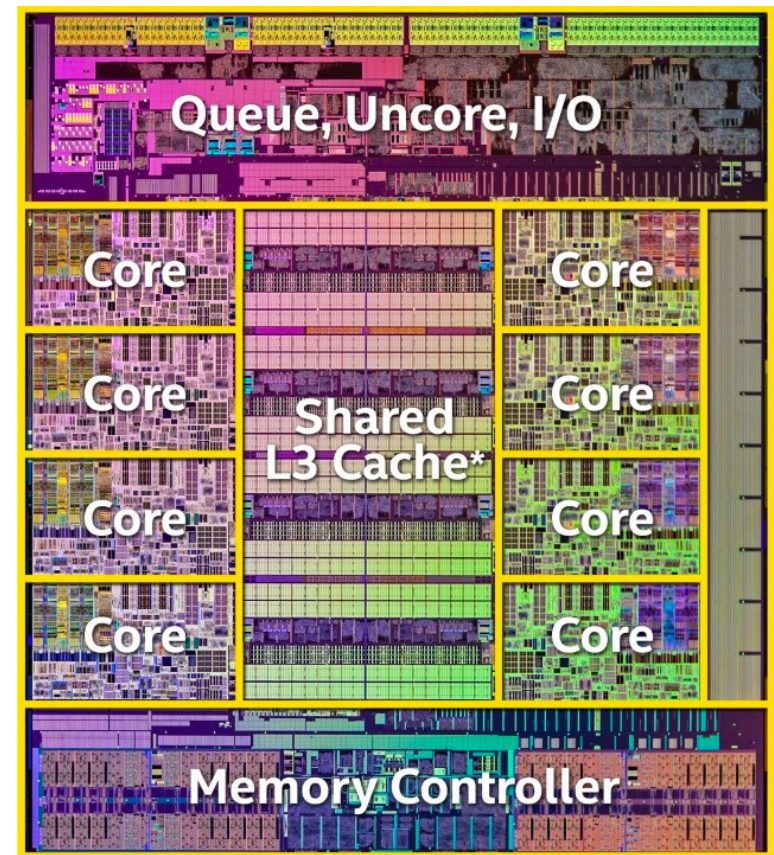
- One hardware thread per socket

This all changed with multi-/many-core

- A processor now has multiple cores
- Each core has its own hardware thread (or even multiple ones → SMT)

To add to the confusion

- OSes will report hardware-threads or cores as processors
- No distinction between sockets (by default)



Die picture of an Intel Xeon Processor

cat /proc/cpuinfo (under Linux)



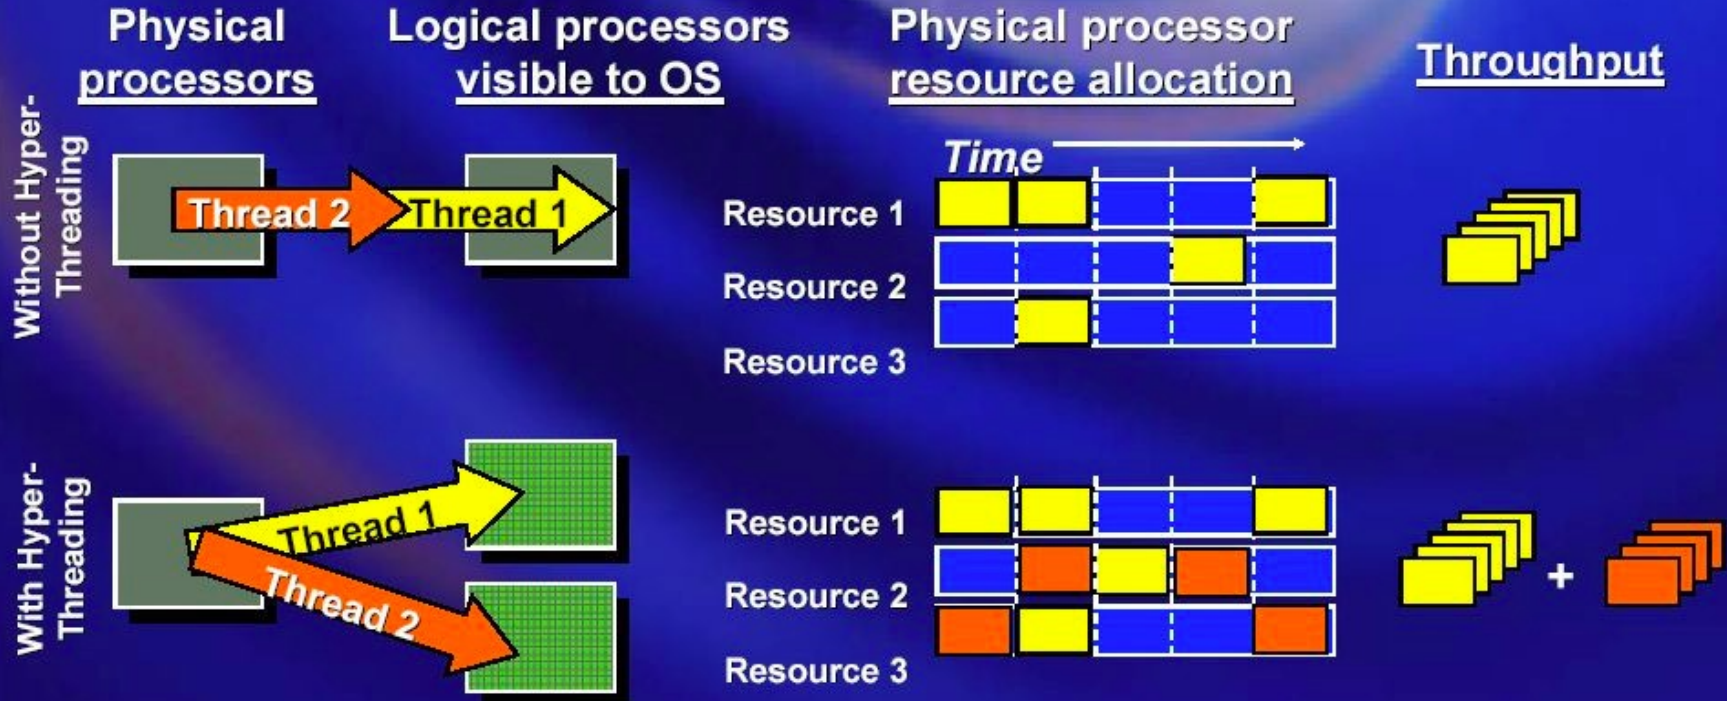
One entry for each “CPU”

i.e., hardware thread

```
processor           : 23
vendor_id           : GenuineIntel
cpu family          : 6
model               : 85
model name          : Intel(R) Xeon(R) Silver 4116
CPU @ 2.10GHz
stepping            : 4
microcode           : 0x2000043
cpu MHz             : 2100.000
cache size          : 16896 KB
physical id        : 1
siblings           : 12
core id            : 13
cpu cores          : 12
apicid              : 58
initial apicid      : 58
fpu                 : yes
fpu_exception       : yes
cpuid level         : 22
wp                  : yes
flags                : fpu vme de pse tsc msr pae
                     : mce cx8 apic sep mtrr pge mca cmov pat pse36
```

```
clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe
syscall nx pdpe1gb rdtscp lm constant_tsc art
arch_perfmon pebs bts rep_good nopl xtopology
nonstop_tsc aperfmperf eagerfpu pni pclmulqdq
dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdbg
fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic
movbe popcnt tsc_deadline_timer aes xsave avx f16c
rdrand lahf_lm abm 3dnowprefetch epb invpcid_single
intel_pt rsb_ctxsw spec_ctrl retpoline kaiser
tpr_shadow vnmi flexpriority ept vpid fsgsbase
tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm
cqm mpx avx512f rdseed adx smap clflushopt clwb
avx512cd xsaveopt xsavec xgetbv1 cqm_llc
cqm_occup_llc cqm_mbm_total cqm_mbm_local
dtherm ida arat pln pts hwp hwp_act_window
hwp_pkg_req
bugs                : cpu_meltdown spectre_v1
spectre_v2
bogomips             : 4191.73
clflush size         : 64
cache_alignment      : 64
address sizes        : 46 bits physical, 48 bits
virtual
power management:
```


How Hyper-Threading Technology Works



Hyper-Threading helps fill moments of idle utilization, such as with:

- Memory accesses (like digital photo editing/effects)
- Dependency chains with longer instruction latencies (like video encoding/transcoding)
- Branch mis-predicts (like 3D ray tracing)
- An integer app and a floating-point app running at the same time

Source: Intel

Using Hyperthreading / SMT

Hyperthreads will also show up as “CPUs”

- BIOS initializes them along with cores and sockets as boot
- Not distinguishable by default
- Changes require reboot

Useful for regular “concurrent” workloads

- E.g., multiple concurrent programs on a laptop
- Resource sharing beneficial

For parallel programming, this is problematic

- Multiple instances of same program with same resource requirements
- Heavy impact on speedup
- Need to be careful on what to schedule on a HT/SMT and what not
- Note: many HPC centers turn this off by default
- Some architectures may require it, though

Still could be useful for background tasks

- System daemons
- I/O operations

What is a Thread?

Independent stream of execution

- Own PC
- Own Stack

Hardware threads

- Implementation of an execution stream in hardware
(think realization of a von Neumann machine in hardware)
- Separate Control Unit executing a sequence of instructions

Software threads

- Programming abstraction that represents a stream of execution
- Seen by programmer

To execute a program

- Programmer defines a software thread
- Software thread gets mapped to hardware thread for execution
(by OS and/or runtime)

Software Threading Basics

Traditional view

- Operating systems maintain processes
- Processes get scheduled to available hardware threads (aka. processors)
- Each process has one execution stream

Processes maintain isolation for protection

- Separate address spaces and files
- Coupled with user IDs
- Communication only via IPC (Inter-Process Communication)

Threading was intended to make this easier

- Sharing of data without protection boundaries
- Cooperative concurrency to support asynchronous behavior
 - Example: I/O in the background, GUI threads
- OS still responsible for scheduling (at least for system-level threads)
 - Enables preemption and progress

Most Common Model: Fork/Join Model

A main thread executes in a process

- Sequential execution on a HW thread

Process want to spawn a second thread

- E.g., to start handing off work

“Forks” a new thread


- Starts a new thread
- New threads start at a specified routine

Both threads operate concurrently

- One or more hardware threads
- Location (at first) transparent

At the end of the parallel regions:

- Main thread waits for other thread to complete
- Continues sequential execution



Different APIs and Threading Implementations

How to implement the fork operation?

- How can the fork be specified by the programmer?
- Is this an OS operation?
 - Mapping to OS abstractions
- Is this a runtime (user-level) runtime mechanism?
 - Understanding thread properties

Similar for join operation

How to pass parameters into a new thread?

How to retrieve results?

- Arguments for new computation
- Results of computation

Runtime interaction

- Synchronization between threads while they are running

The Linux Clone call (System Threads)

```
int clone(int (*fn)(void *), void *child_stack, int flags,  
void *arg, ... /* pid_t *ptid, void *newtls, pid_t *ctid */ );
```

Shared call for process and thread creation

Flags	CLONE_FILES	Parent/Child share file descriptor table
	CLONE_NEWIPC	Establish new IPC name space
	CLONE_NEWNET	Establish new network name space
	CLONE_NEWUSER	Create new child under new UID
	CLONE_THREAD	Parent/Child in same thread group
	CLONE_VM	Parent/Child in same memory space
	...	

Threads and processes are scheduled by the same scheduler

Arguments passed via “arg”, return via function return

User-Level Threads

Alternative: implement threads as part of a user-level library

- Maintain own PC and stack
- Switch between threads as needed
- No kernel support or modifications necessary

Advantages

- Easier to implement and support
- Lighter weight


Disadvantages

- No scheduling guarantees
- Preemption and progress is very hard to guarantee, if not impossible
- Often requires explicit `yield()` calls required

Use cases

- Light-weight, fine-grained task based parallel programming
- Closely coordinated activities with well-define switchpoints

Hybrid Models (e.g., Solaris)



From: Brian Vinter, University of Copenhagen, <http://www.diku.dk/~vinter/xmp/lecture2.pdf>

The Windows View of the World

A process is a data structure with

- Memory space
- File descriptors
- Network interfaces
- BUT: a process cannot execute on its own

A process contains one more threads

- Threads have a stack and a program counters
- Threads execute

Additional: Fibers

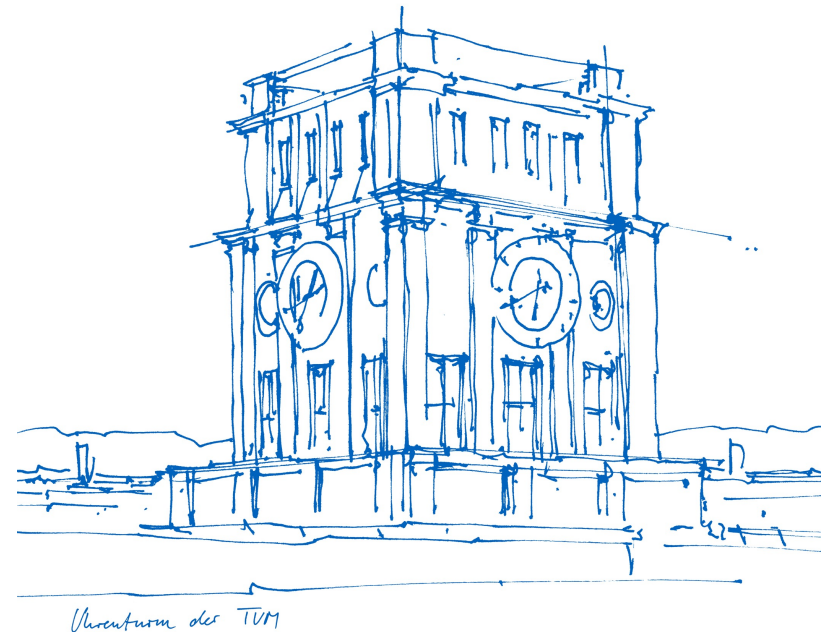
- User-level threads
- User scheduled and cannot be preempted
- Hierarchical compared to main threads

Lecture IN-2147

Parallel Programming

Martin Schulz
Technical University of Munich
Department of Computer Engineering

Segment 8
The POSIX Thread API



POSIX Threading

POSIX: Portable Operating System Interface

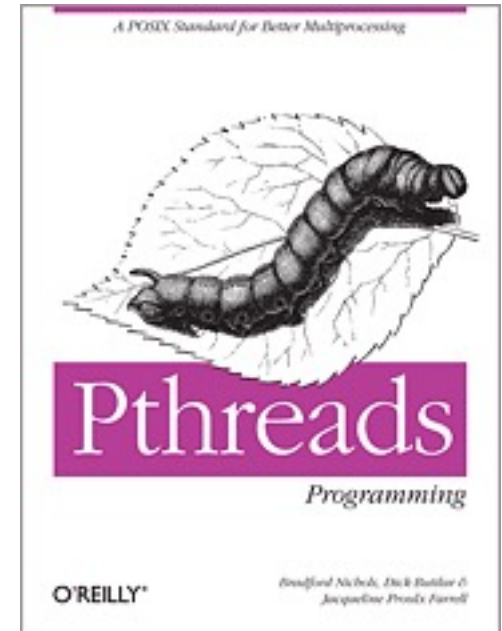
- Family of IEEE Standards
- Ensure compatibility between OSs
- Defines and API, Shells, Utilities

POSIX threads: IEEE Std 1003.1c-1995

- Contains about 100 procedures (Prefix `pthread_`)
- Standardized access to threads
 - Creation, destruction
 - Coordination and synchronization
 - Thread management
- Available on most systems
- Typically used for system level threads

Programming with pthreads

- `#include <pthread.h>`
- On some systems compile with „-lpthreads“



By Dick Buttlar, Jacqueline
Farrell, Bradford Nichols
Publisher: O'Reilly Media
Release Date: January 2013

Also:

[https://computing.llnl.gov/
tutorials/pthreads/](https://computing.llnl.gov/tutorials/pthreads/)

POSIX Thread Create

```
int pthread_create(pthread_t *thread,  
const pthread_attr_t *attr,  
void *(*start_routine) (void *), void *arg);
```

Create a new Pthread

Returns 0 if successful

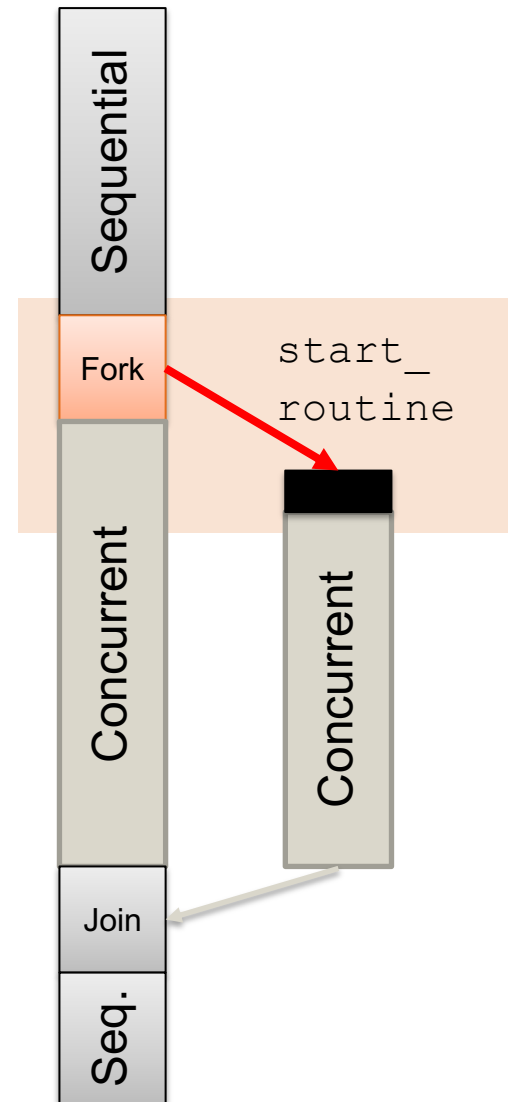
Arguments:

thread = returns thread ID

attr = attributes for the new thread

start_routine = routine that executes the new thread

arg = argument passed to the new thread



POSIX Thread Join

```
int pthread_join(pthread_t thread,  
void **retval);
```

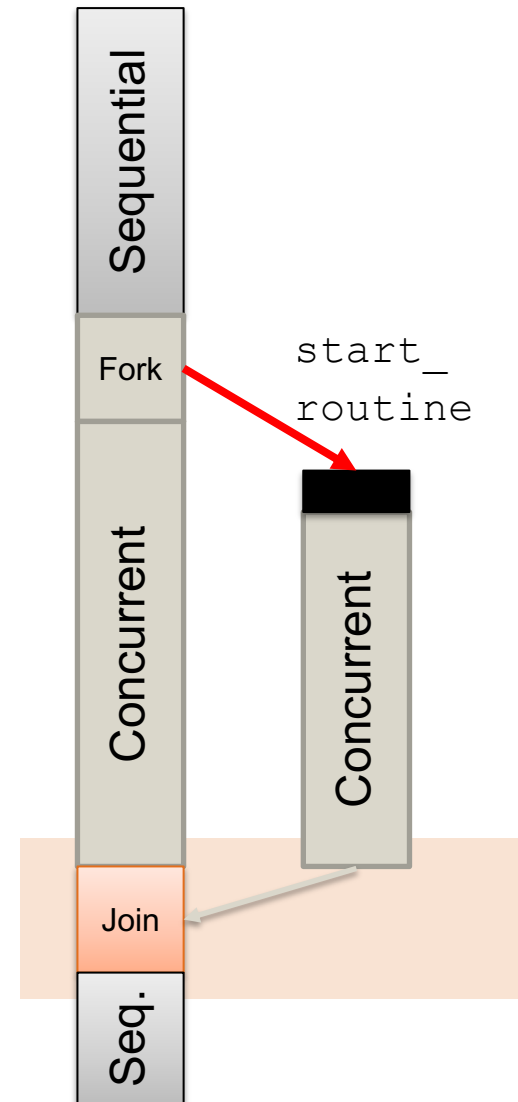
Waits for another thread to terminate

Returns 0 if successful

Arguments:

thread = ID of thread to wait for

retval = return value from the terminated thread



POSIX Threads - Details

Attributes

- Set of properties defining thread behavior
- Examples: bound/unbound, scheduling policy, ...



Thread management and special routines:

- Get own thread ID: `pthread_self()`
- Compare two thread IDs: `pthread_equal(t1, t2)`
- Run a particular function once in a process: `pthread_once(ctrl, fct)`

Stack management

- Routines to set and get the stack size
- Routines to set and get the stack address

Stack Structure



Fork/Join Example (part 1)

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#define NUM_THREADS 5

void* perform_work(void* argument)
{
    int passed_in_value;
    passed_in_value = *((int*) argument);
    printf("Hello World! It's me, thread with argument
           %d!\n", passed_in_value);
    return NULL;
}
```


Fork/Join Example (part 2)

```
int main(int argc, char** argv)
{
    pthread_t threads[NUM_THREADS];
    int thread_args[NUM_THREADS];
    int result_code; unsigned index; /

    / create all threads one by one
    for (index = 0; index < NUM_THREADS; ++index)
    {
        thread_args[ index ] = index;
        printf("In main: creating thread %d\n", index);
        result_code = pthread_create(&threads[index],
                                     NULL,
                                     perform_work, &thread_args[index]);
        assert(!result_code);
    }
}
```

Source: Wikipedia

Fork/Join Example (part 3)

```
// wait for each thread to complete
for (index = 0; index < NUM_THREADS; ++index)
{
    // block until thread 'index' completes
    result_code = pthread_join(threads[index], NULL);
    assert(!result_code);
    printf("In main: thread %d has completed\n",
        index);
}
printf("In main: All threads completed successfully\n");

exit(EXIT_SUCCESS);

}
```

Fork/Join Example (Output)

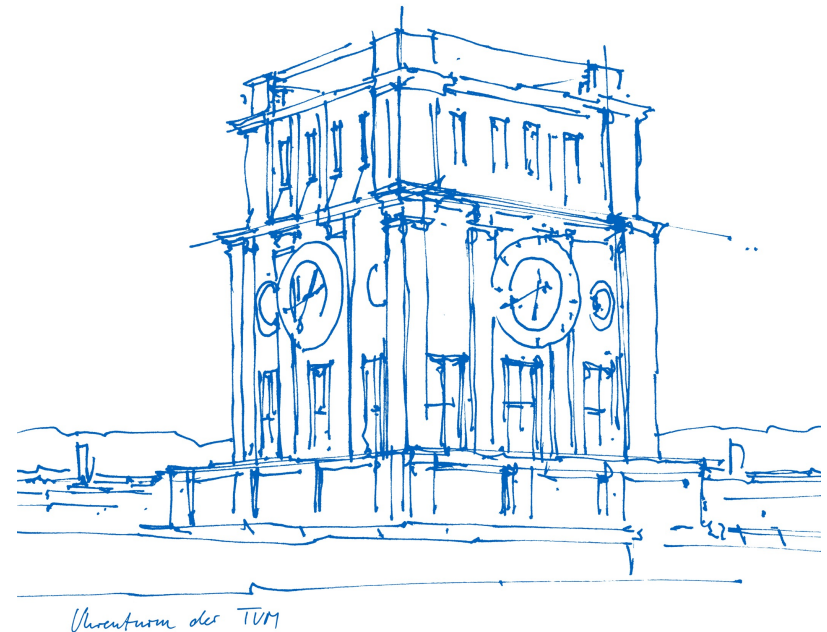
```
In main: creating thread 0
In main: creating thread 1
Hello World! It's me, thread with argument 0!
In main: creating thread 2
Hello World! It's me, thread with argument 1!
In main: creating thread 3
Hello World! It's me, thread with argument 2!
In main: creating thread 4
Hello World! It's me, thread with argument 3!
Hello World! It's me, thread with argument 4!
In main: thread 0 has completed
In main: thread 1 has completed
In main: thread 2 has completed
In main: thread 3 has completed
In main: thread 4 has completed
In main: All threads completed successfully
```

Lecture IN-2147

Parallel Programming

Martin Schulz
Technical University of Munich
Department of Computer Engineering

Segment 9
Thread Synchronization



Synchronization Between Threads

Unless we deal with pure concurrency

- Independent tasks
- No dependencies

we need to coordinate between threads

Examples

- Enforce common completion of tasks
- Enforce happens before relationships
- Guard updates to common data structures

In POSIX threads, there are two main concepts

- Locks / Mutual Exclusion (passive waiting for resources)
- Condition variables (active signaling, covered in background material)

Other constructs can be built on top of them

Locks for Mutual Exclusion

Problem: concurrent access to shared resources

- Shared variables, memory locations
- Access to I/O
- Two or more threads concurrent updates can lead to inconsistencies

Classic example: depositing money into a bank account

```
int account = 100;
void deposit(int money)
{
    account = account + money;
}
```

Thread Alice: `deposit(200);`

Thread Bob: `deposit(100);`

Possible final values for account:

400

both succeed

300

both read
Bob writes first

200

both read
Alice writes first

POSIX Thread Locks

Initialization of a lock:

- Global variable of type `pthread_mutex_t`
- Initialize to `PTHREAD_MUTEX_INITIALIZER`
- Can also be done dynamically: `pthread_mutex_init/destroy()`

Lock a mutex

- `pthread_mutex_lock(&mutex);`
- Blocks until mutex is granted

Unlock a mutex

- `pthread_mutex_unlock(&mutex);`
- Returns immediately

Lock a mutex, if available

- `pthread_mutex_trylock(&mutex);`
- Returns immediately

Also locks can have attributes

POSIX Lock Example

```
int account = 100;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

void deposit(int money)
{
    pthread_mutex_lock(&mutex);
    account = account + money;
    pthread_mutex_unlock(&mutex);
    return 0;
}
```

Note:

- mutex is a standalone variable
- Not explicitly associated with the memory it protects
- User/programmer responsibility

Lock Implementations

Criteria for implementations

- Correctness: Guarantees mutual exclusion
- Fairness: Every process eventually gets the mutex (and on average “equally fast”)

Spin-Locks

- If lock is taken, actively wait by “spinning” on a flag
- Implementation typically via atomic operations
- Advantage: fast response time
- Disadvantage: blocks the hardware threads, uses resources, costs energy
- Generally bad for concurrent executions

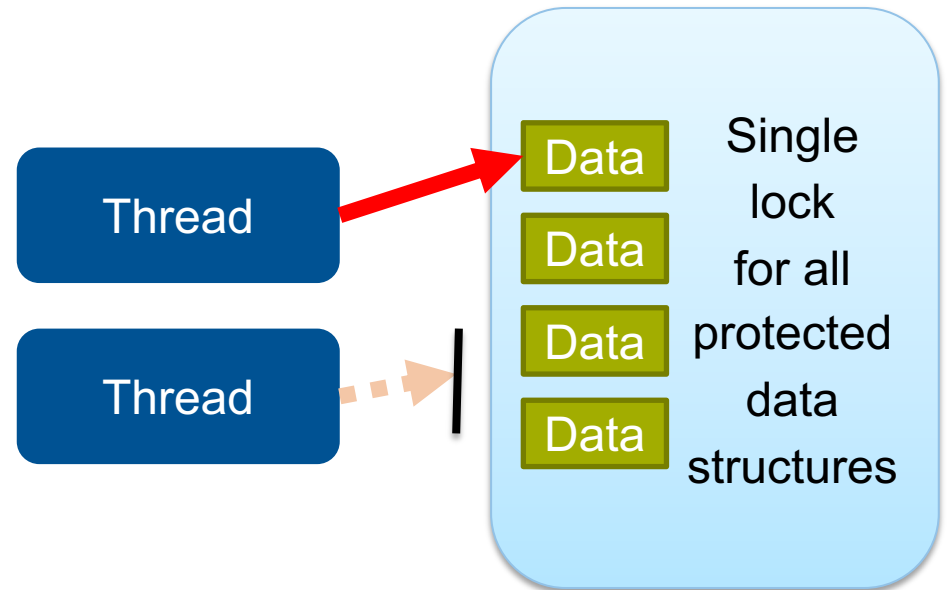
Yielding Locks

- If lock is taken, yield hardware thread
- Implementation typically via runtime system
- Advantage: low resource usage
- Disadvantage: slow response time
- Generally bad for HPC

Lock Granularity

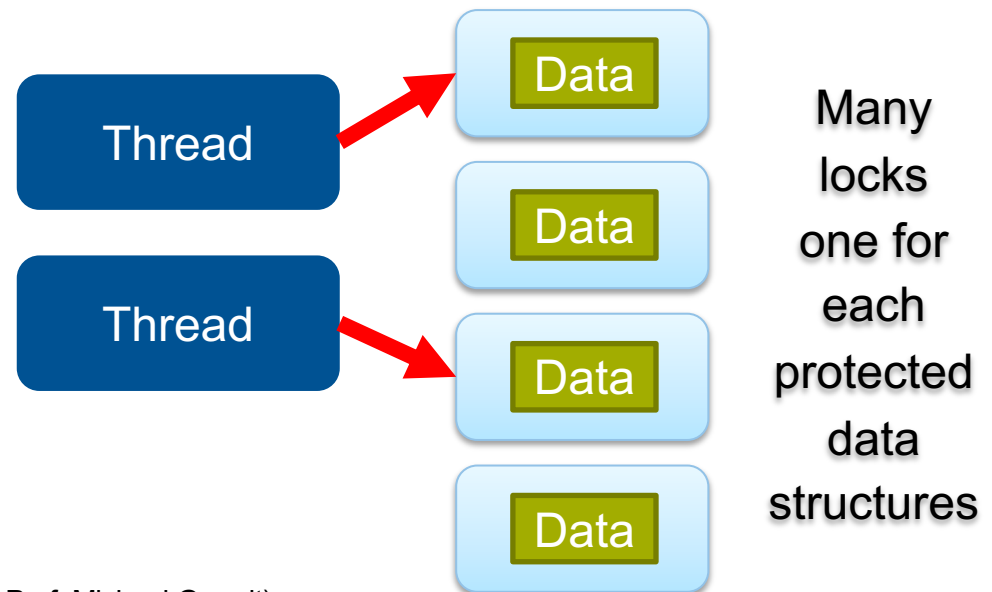
Coarse grained locking

- One single lock for all data
- Limits concurrency
- Eases implementations
- Older OSes use this



Fine grained locking

- One lock for each data element
- Maximizes concurrency (multiple threads can have locks)
- Requires many locking calls
- May require multiple locks at the same time

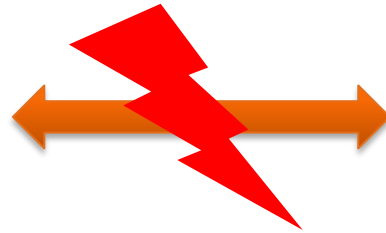


Hybrid versions often useful

Danger: Deadlocks

Thread 1

```
LOCK FOR A
LOCK FOR B
READ A and B
UNLOCK A
UNLOCK B
```

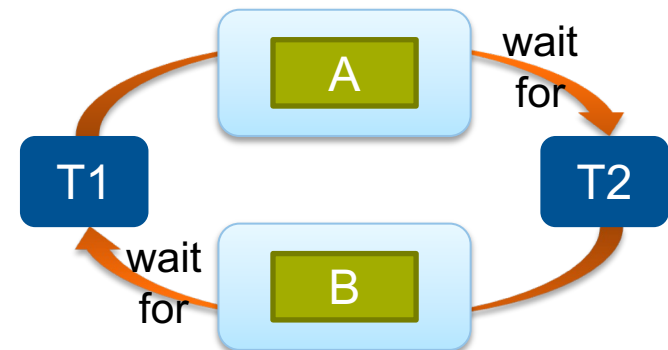


Thread 2

```
LOCK FOR B
LOCK FOR A
READ A and B
UNLOCK A
UNLOCK B
```

Both threads request lock A and B

- Both end up waiting for the second lock
- Cyclic dependency!



Classic Example: Dining Philosophers

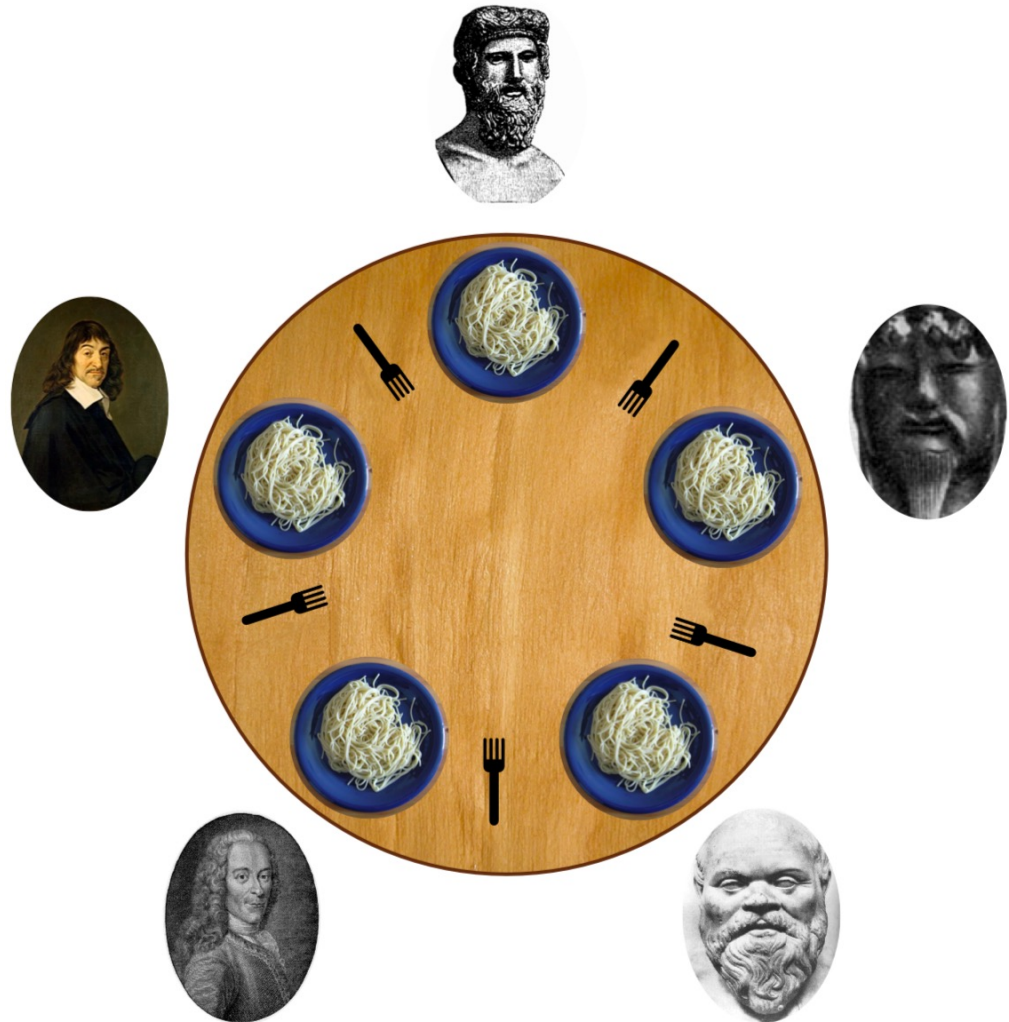
5 philosophers sit at a table and

- Think
- Think until left fork is free
- Pick up left fork
- Think until right fork is free
- Pick up right fork
- Eat for some time
- Put down both forks

Translated to locks:

Philosopher P

- Compute
- `Lock(lock[P])`
- `Lock(lock[(P+1) % 5])`
- Compute
- `Unlock(lock[P])`
- `Unlock(lock[(P+1) % 5])`



Source: Wikipedia

Strategies for Deadlock Avoidance

Easy option: only hold one lock at a time

Central arbiter to ask for permission

Order among all locks and only allocate in that order

Thread 1

LOCK FOR A

LOCK FOR B

READ A and B

UNLOCK A

UNLOCK B



Thread 2

LOCK FOR A

LOCK FOR B

READ A and B

UNLOCK A

UNLOCK B

Classic Example: Dining Philosophers

All solutions have drawbacks

One lock

- We still need two forks
- One lock for whole table?

Arbiter

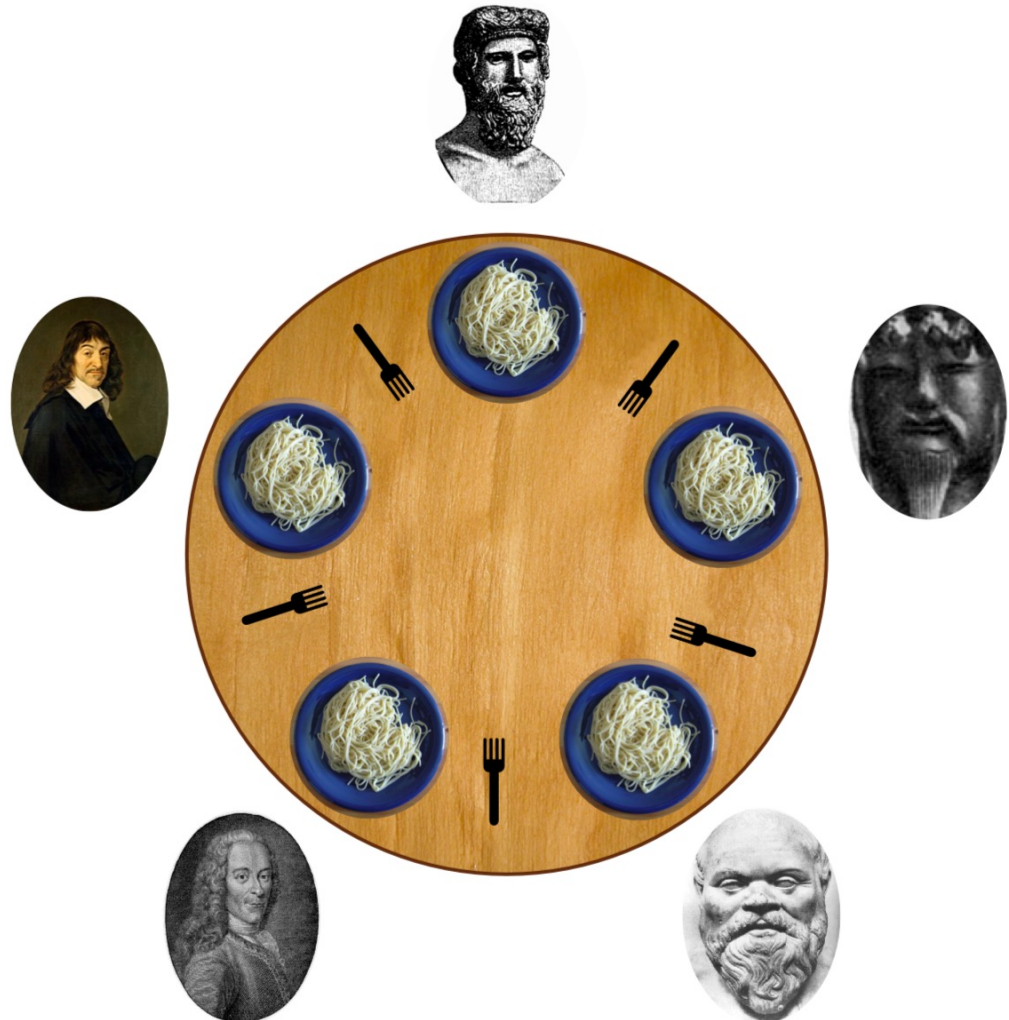
- Need central instance
- Might be complex to implement

Lock order

- Fairness
(one philosopher may starve)

Custom schemes possible

- Requesting forks from neighbors
- Preference to starving processes



Source: Wikipedia

Performance Aspects for Threading

Overheads

- Thread creation and destruction can be expensive operations
- Ensure large parallel regions or “park” threads

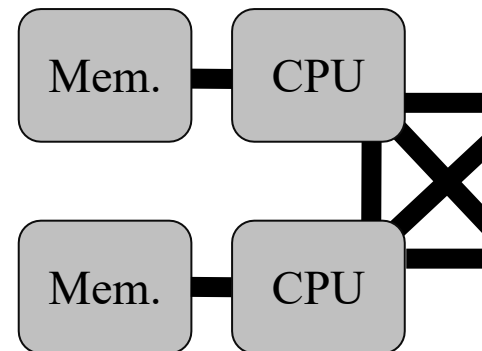
Lock contention

- Locks are low-overhead when few threads try access them
- Overhead grows with more threads accessing them more often

Thread pinning

- For system-level threads the OS does scheduling
 - Mapping of SW to HW thread
 - Determines the location of a thread's execution
- Large impact on performance
 - Determines what is needed to share information
 - NUMA properties
- Pinning, fixing a SW thread to a (group of) HW thread(s)
 - Thread attributes, libraries like libNUMA (see `man numa`)

Good use of caches, avoidance of false sharing

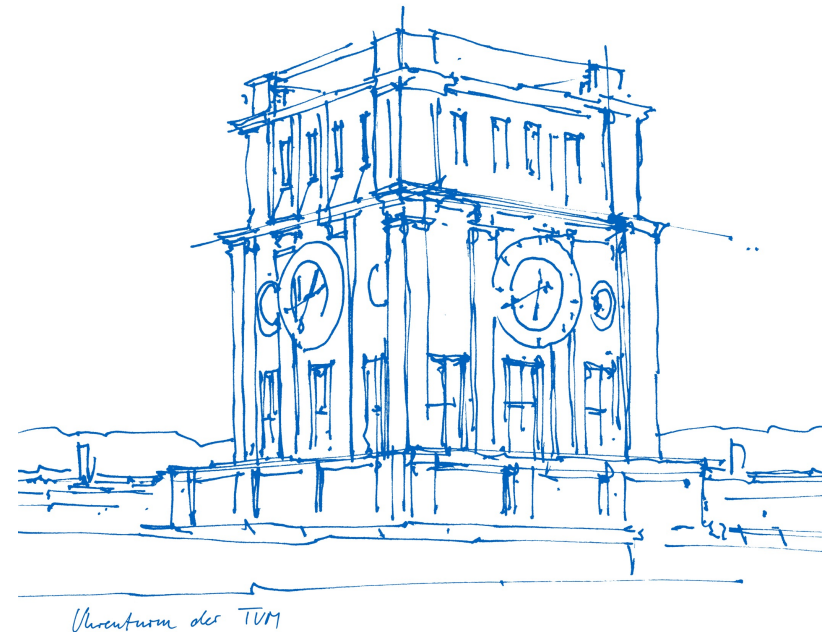


Lecture IN-2147

Parallel Programming

Martin Schulz
Technical University of Munich
Department of Computer Engineering

Segment 10
Modern Thread APIs



Different Thread APIs

Native thread APIs per system

- POSIX Threads
- Win32 Threads
- Solaris Threads

Portable standards

- OpenMP

Many custom/research packages

- Often mapped to native APIs
- Often user-level threads

Motivation for custom APIs

- Lower overhead
- Customized for particular tasks
- Custom hardware with special properties

Native thread APIs per language

- C++ threads
- Java threads
- Rust threads

(requires a runtime system)

Language APIs can Simplify Usage (e.g., C++)

```
#include <string>
#include <iostream>
#include <thread>
using namespace std;
```

Task to be run
in thread

```
// The function we want to execute on the new thread.
void task1(string msg)
{ cout << "task1 says: " << msg; }
```

```
int main()
{
    // Constructs the new thread and runs it. Does not block.
    thread t1(task1, "Hello");
```

Fork

```
// Do other things...
```

Join

```
// Join threads, blocks until completion
t1.join();
```

```
}
```

On G++, compile with `-std=c++0x -pthread`.

Source: <https://stackoverflow.com/questions/266168/simple-example-of-threading-in-c>

Mutex Example C++

```
#include <iostream>
#include <thread>
#include <mutex>
```

Declaration

```
std::mutex mtx;
```

```
void print_thread_id (int id)
{
    mtx.lock();
    std::cout <<
        "thread #" << id << '\n';
    mtx.unlock();
}
```

Lock

Unlock

```
int main ()
{
    std::thread threads[10];

    // spawn 10 threads
    for (int i=0; i<10; ++i)
        threads[i] =
            std::thread(
                print_thread_id, i+1);

    for (auto& th : threads)
        th.join();

    return 0;
}
```

Source: <https://www.cplusplus.com/reference/mutex/mutex/lock/>

Example: Java threads

```
public class ThreadExample
```

```
{
```

```
    public static void main(String[] args)
```

```
    {
```

```
        System.out.println(Thread.currentThread().getName());
```

```
        for(int i=0; i<10; i++)
```

```
        {
```

```
            new Thread("" + i)
```

```
            {
```

```
                public void run()
```

```
                {
```

```
                    System.out.println("Thread: " + getName() + " running");
```

```
                }
```

```
            }.start();
```

```
        }  
    }  
}
```

New Thread Declaration

Task to be run
in thread (class)

Fork / Thread Start

Class also has a "join" method

Example: Java Locks

```
public class Counter
```

```
{
```

```
    private Lock lock = new Lock();
```

```
    private int count = 0;
```

Declaration

```
    public int inc()
```

```
    {
```

```
        lock.lock();
```

Lock

```
        int newCount = ++count;
```

```
        lock.unlock();
```

```
        return newCount;
```

```
    }
```

```
}
```

Unlock

Summary

Many thread libraries follow the fork/join model

- Mapped onto language constructs
- Very similar functionality
- Costs may vary widely, though, requiring different tradeoffs

Mutual exclusions is also a very basic primitive

- Given in many models for parallelism and concurrency
- Typical multiple options for locks (try-locks, multi-locks, etc.)

Often additional synchronization mechanisms

- Signaling like in condition variables
- Explicit dependencies

Implementation at the end based on low-level threads

- User level threads in the runtime
- OS threads or abstracts like Pthreads



Exercises
with
C++
threads