

# Lecture IN-2147

## Parallel Programming


Martin Schulz

Technical University of Munich

Department of Computer Engineering

### Lecture 3: OpenMP 101

- Seg. 11: OpenMP and its Execution Model
- Seg. 12: OpenMP Syntax and Basics
- Seg. 13: Data Sharing in OpenMP
- Seg. 14: Synchronization in OpenMP



# Lecture IN-2147

## Parallel Programming


Martin Schulz

Technical University of Munich

Department of Computer Engineering

### Lecture 3: OpenMP 101

- Performance Aspects for Threading
- Seg. 10: Other Thread APIs
- Seg. 11: OpenMP and its Execution Model
- Seg. 12: OpenMP Syntax and Basics
- Seg. 13: Data Sharing in OpenMP
- Seg. 14: Synchronization in OpenMP



# Performance Aspects for Threading

## Overheads

- Thread creation and destruction can be expensive operations
- Ensure large parallel regions or “park” threads


## Lock contention

- Locks are low-overhead when few threads try access them
- Overhead grows with more threads accessing them more often

## Thread pinning

- For system-level threads the OS does scheduling
  - Mapping of SW to HW thread
  - Determines the location of a thread's execution
- Large impact on performance
  - Determines what is needed to share information
  - NUMA properties
- Pinning, fixing a SW thread to a (group of) HW thread(s)
  - Thread attributes, libraries like libNUMA (see `man numa`)

Good use of caches, avoidance of false sharing




# Lecture IN-2147

## Parallel Programming

Martin Schulz  
Technical University of Munich  
Department of Computer Engineering

Segment 10  
Modern Thread APIs



# Different Thread APIs

## Native thread APIs per system

- POSIX Threads
- Win32 Threads
- Solaris Threads

## Portable standards

- OpenMP

## Many custom/research packages

- Often mapped to native APIs
- Often user-level threads

## Motivation for custom APIs

- Lower overhead
- Customized for particular tasks
- Custom hardware with special properties

## Native thread APIs per language

- C++ threads
- Java threads
- Rust threads

(requires a runtime system)

# Language APIs can Simplify Usage (e.g., C++)

```
#include <string>
#include <iostream>
#include <thread>
using namespace std;
```

Task to be run  
in thread

```
// The function we want to execute on the new thread.
void task1(string msg)
{ cout << "task1 says: " << msg; }
```

```
int main()
{
    // Constructs the new thread and runs it. Does not block.
    thread t1(task1, "Hello");
```

Fork

```
// Do other things...
```

Join

```
// Join threads, blocks until completion
t1.join();
```

```
}
```

*On G++, compile with `-std=c++0x -pthread`.*

Source: <https://stackoverflow.com/questions/266168/simple-example-of-threading-in-c>

# Mutex Example C++

```
#include <iostream>
#include <thread>
#include <mutex>
```

Declaration

```
std::mutex mtx;
```

```
void print_thread_id (int id)
{
    mtx.lock();
    std::cout <<
        "thread #" << id << '\n';
    mtx.unlock();
}
```

Lock

Unlock

```
int main ()
{
    std::thread threads[10];

    // spawn 10 threads
    for (int i=0; i<10; ++i)
        threads[i] =
            std::thread(
                print_thread_id, i+1);

    for (auto& th : threads)
        th.join();

    return 0;
}
```

Source: <https://www.cplusplus.com/reference/mutex/mutex/lock/>

# Example: Java threads

```
public class ThreadExample
```

```
{
```

```
    public static void main(String[] args)
```

```
    {
```

```
        System.out.println(Thread.currentThread().getName());
```

```
        for(int i=0; i<10; i++)
```

```
        {
```

```
            new Thread("" + i)
```

```
            {
```

```
                public void run()
```

```
                {
```

```
                    System.out.println("Thread: " + getName() + " running");
```

```
                }
```

```
            }.start();
```

```
        }  
    }  
}
```

New Thread Declaration

Task to be run  
in thread (class)

Fork / Thread Start

Class also has a "join" method



# Example: Java Locks

```
public class Counter
```

```
{
```

```
    private Lock lock = new Lock();
```

```
    private int count = 0;
```

Declaration

```
    public int inc()
```

```
{
```

```
    lock.lock();
```

```
    int newCount = ++count;
```

```
    lock.unlock();
```

```
    return newCount;
```

```
}
```

```
}
```

Lock

Unlock

# Recap: Lecture 2 - Threading

A thread is a stream of execution

Hardware threads

- Implementation of the “von Neumann” control unit (CU)
- Complicated by multi-/many core and Hyperthreading/SMT

Software threads are abstracting execution streams for the programmer

- Distinguish user- and system-level threads
- Hybrid systems exist in some operating systems
- Real concurrency has to be layered on top of system-level threads

Thread execution by mapping SW thread to HW threads

- HW thread will execute SW thread
- Mapping (which core, socket, ...) important for performance

Widely available and standardized API: POSIX threads

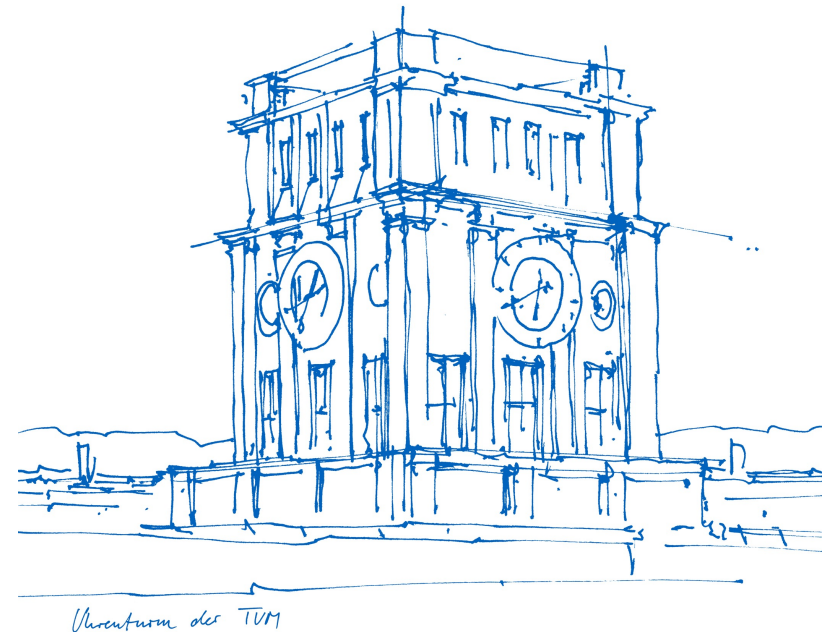
- Routines for thread creation/destruction
- Several synchronization constructs, incl. mutexes and condition variables

# Lecture IN-2147

## Parallel Programming

Martin Schulz  
Technical University of Munich  
Department of Computer Engineering

Segment 11  
OpenMP and its Execution Model



# Pros and Cons of Using Pthreads

Pthreads represent direct abstraction of system-level parallelism

- Direct map to underlying hardware threads (in most cases)
- Even parallel loops are difficult to think about

Association of data and threads is up to the programmer

- No program construct to do this
- Data locality has to be done explicitly

Need to manage data visibility by hand

- Which variables are per thread?
- Which variables are shared across all threads?

Simple synchronization primitives

- Locks and condition variables
- Not sufficient to do easy parallel coordination

# Finding a Higher Level of Abstraction

Initially, each OS had its own threading abstraction

- All low-level and library based
- Even with the standard Pthreads,

In the 1990's several vendors supplied their own language extensions

- Driven by rising popularity of shared memory machines
- Mainly as extensions to Fortran
- Users could specify parallel loops
- A compiler would then translate the code to a threaded program

Very helpful for users

- Small, often incremental additions to code
  - In many cases, one line or one pragma
  - Codes stayed readable
- Compiler/runtime would do the engineering work

But: each vendor had their extensions, which made codes non-portable

- First attempt at standardizing failed (ANSI X3H5) in 1994
- A few years later, another attempt led to OpenMP

# OpenMP = Open Multi-Processing

Standard for writing parallel programs, mainly “on-node”

- Standardization across many vendors, systems, architectures, ...
- “Lean and Mean” – simple API to achieve complex goals
- Ease of Use
- Portability

Managed by the OpenMP ARB

- Architecture Review Board
- Independent non-profit organization
- Members are companies, universities, research labs

First version published in 1997

- Many advances since then
  - Tasking, SIMD, GPU support
- Currently at Version 5.2, published in November 2021
- All documents (and more) at [www.openmp.org](http://www.openmp.org)

# OpenMP

OpenMP is ...

... an Application Program Interface (API) that may be used to program multi-threaded, shared memory parallelism (plus accelerators)

... comprised of three primary API components:

- Compiler Directives (for C/C++ and Fortran)
- Runtime Library Routines
- Environment Variables

OpenMP is not ...

- ... intended for distributed memory systems
- ... necessarily implemented identically by all vendors
- ... guaranteed to automatically make the most efficient use of shared memory
- ... required to check for data dependencies, race conditions, deadlocks, etc.
- ... designed to handle parallel I/O

# A Simple Example

```
#include <omp.h>

main() {
    #pragma omp parallel
    {
        printf("Hello world");
    }
}
```

## Compilation

- `icc -O3 -openmp openmp.c`
- `gcc -O3 -fopenmp openmp.c`

- This differs between compilers

```
> export OMP_NUM_THREADS=2
> a.out
Hello world
Hello world

> export OMP_NUM_THREADS=3
> a.out
Hello world
Hello world
Hello world
```




# Fork/Join Execution Model

Older OpenMP Versions  
refer to a “master” thread  
(new deprecated)

## Parallel Regions

1. An OpenMP-program starts as a single thread (*primary thread*).
2. Additional threads are created when the master hits a parallel region.
3. After the parallel region, the new threads are given back to the runtime
4. The *primary thread* continues after the parallel region.

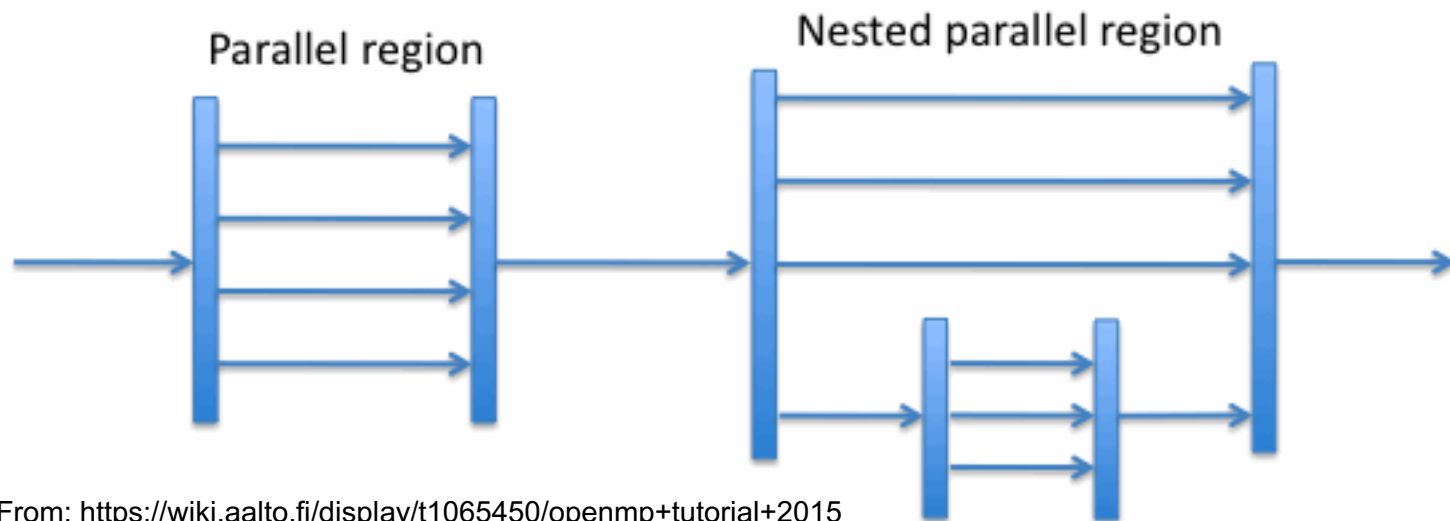
All threads in the team are *synchronized* at the end of a parallel region via a barrier.



Source: LLNL tutorial by Blaise Barney, adapted to OpenMP 5.2

# Nested Parallelism

OpenMP threads can themselves create parallel regions



From: <https://wiki.aalto.fi/display/t1065450/openmp+tutorial+2015>

## Noteworthy

- OpenMP is not required to spawn/use more threads in the nested region
  - It is compliant to just continue executing sequentially
- Mapping to hardware threads is up to the runtime

# OpenMP Implementations

OpenMP is a language extension

- On top of C/C++ or Fortran
- Pragmas to the base language, which can be ignored

Consequence: need a new compiler

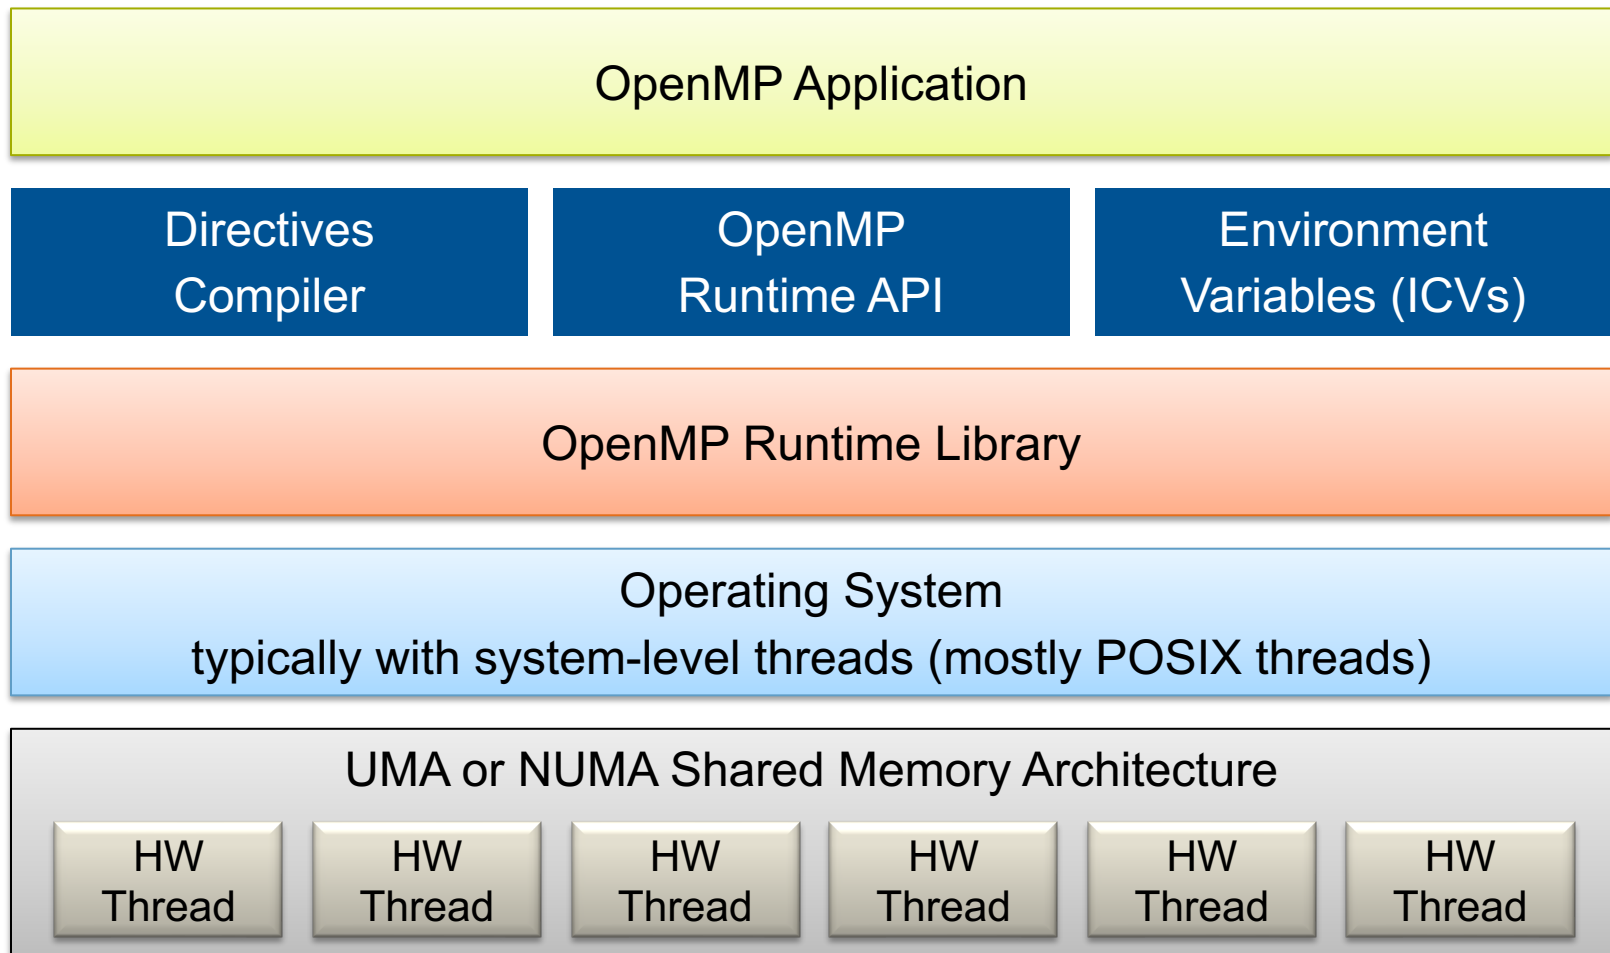
- Implemented within existing compilers
- Most compilers support this now: gcc, icc, LLVM, PGI, ...

Additionally: need a runtime system

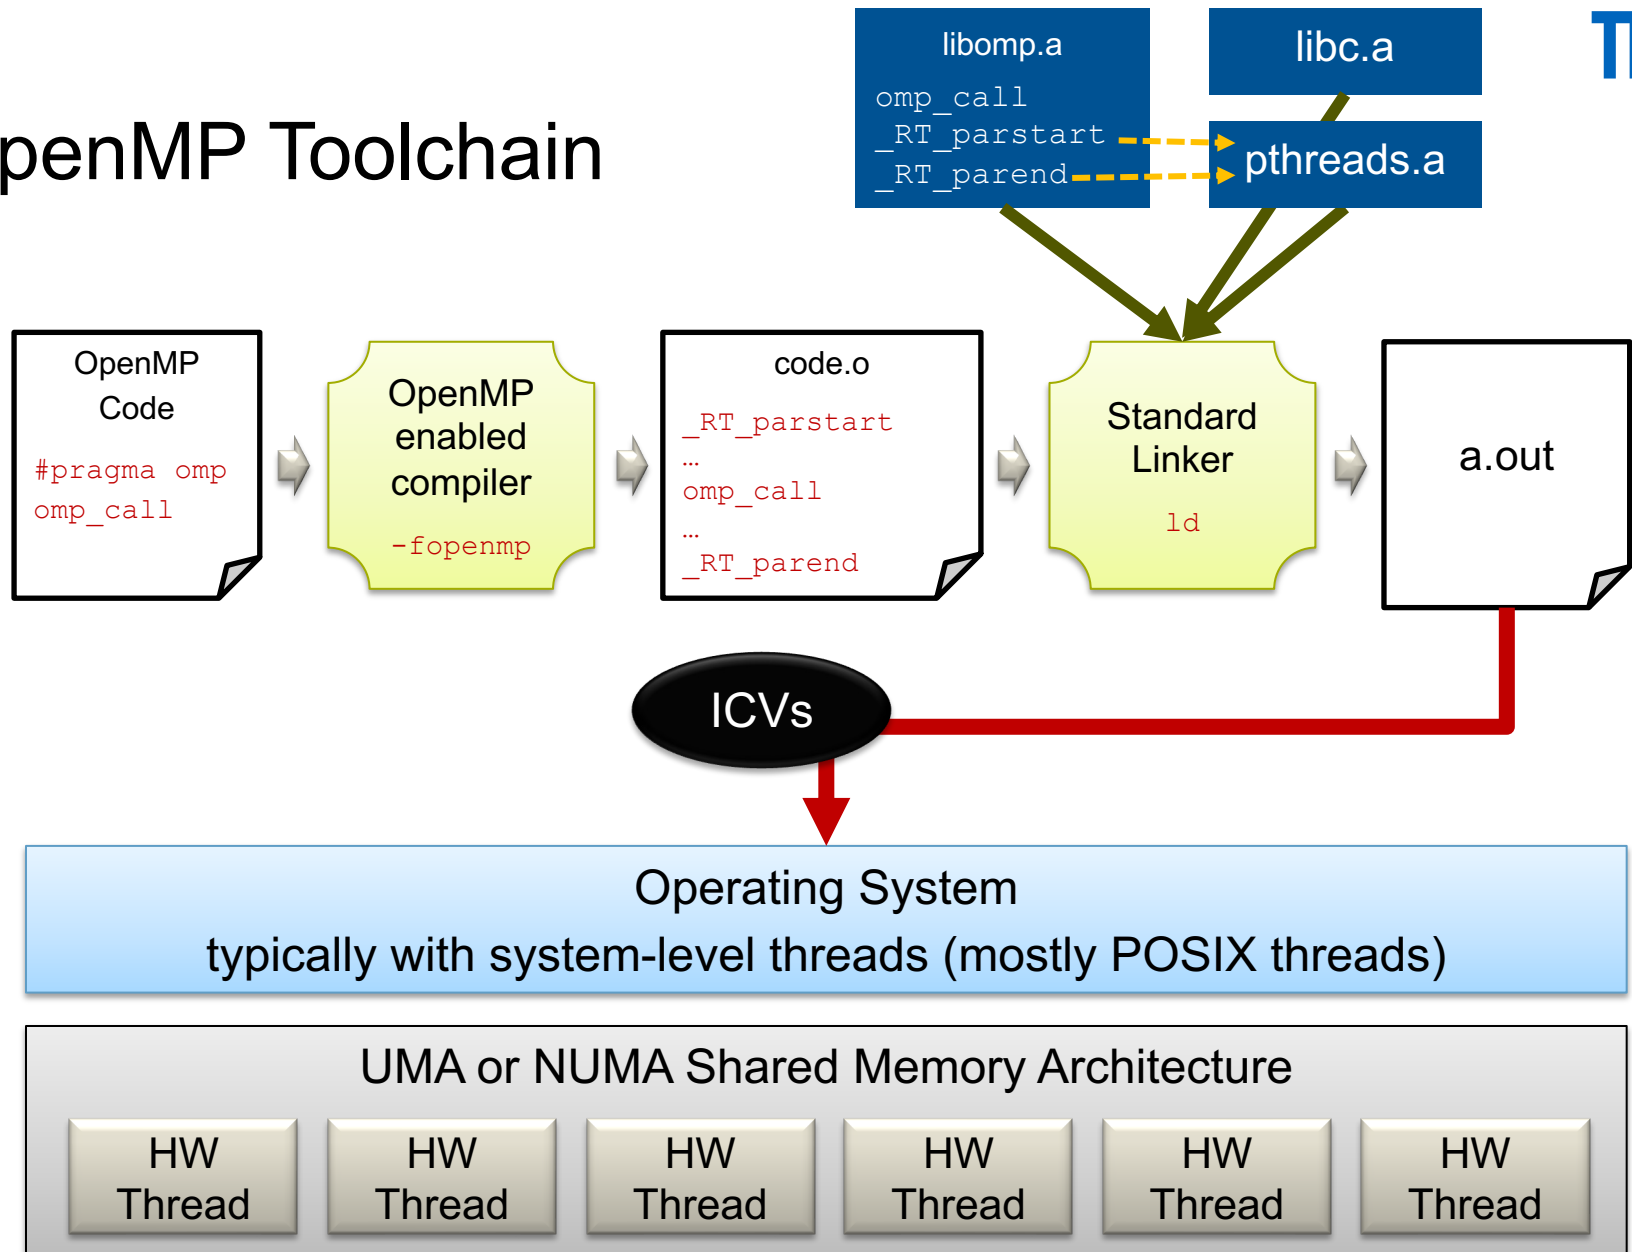
- Maps program to underlying thread package
- Often built on top of Pthreads
- Standard defines some user functions, but NOT the entire interface
- Two widely known open source systems: gomp and Intel's OpenMP/LLVM runtime

Well defined environment variables, also called „internal control variables“ or „ICVs“

# OpenMP System Stack



# OpenMP Toolchain




# Lecture IN-2147

## Parallel Programming

Martin Schulz  
Technical University of Munich  
Department of Computer Engineering

Segment 12 OpenMP Syntax and Basics



# OpenMP Syntax

Most of the constructs in OpenMP are compiler directives

```
#pragma omp construct [clause [clause]...]
```

Example: `#pragma omp parallel`

Most OpenMP constructs apply to a “structured block”.

- A block of one or more statements
- With one point of entry at the top and one point of exit at the bottom
- Think: C/C++

```
{ ... }
```

Additionally: runtime API

- Function prototypes and types in the `omp.h` header file
- Namespace: `omp_`

# Fortran

OpenMP also defines pragmas for Fortran

- Same concepts and similar pragmas
- Different syntax
- Need to deal with scoping differently

**!\$OMP *directive name* [*parameters*]**

Example:

```
!$OMP PARALLEL DEFAULT(SHARED)  
    write(*,*) 'Hello world'  
!$OMP END PARALLEL
```



# More on Directives Syntax

Directives can have continuation lines

- C  
**#pragma omp parallel private(i) \**  
**private(j)**
- Fortran  
**!\$OMP *directive name first\_part* &**  
**!\$OMP *continuation\_part***

# Parallel Regions

Parallel regions are marked explicitly

```
#pragma omp parallel [parameters]  
{  
    block  
}
```

Block executed in parallel by a *team of threads*

Degree of parallelism controlled by ICV

**OMP\_NUM\_THREADS**

# A Simple Example (repeated)

```
#include <omp.h>

main() {
    #pragma omp parallel
    {
        printf("Hello world");
    }
}
```

```
> export OMP_NUM_THREADS=2
```

```
> a.out
```

```
Hello world
```

```
Hello world
```

```
> export OMP_NUM_THREADS=3
```

```
> a.out
```

```
Hello world
```

```
Hello world
```

```
Hello world
```

## Compilation

➤ `icc -O3 -openmp openmp.c`

➤ `gcc -O3 -fopenmp openmp.c`

- This differs between compilers

# Sections in a Parallel Region

Parallel regions can have separate sections

```
#pragma omp sections [parameters]  
{  
    #pragma omp section  
        { ... block ... }  
    #pragma omp section  
        { ... block ... }  
    ...  
}
```

Most “thread like” option

- Must be within a parallel region
- Each section within a sections block is executed once by one thread
- Threads that finished their section wait at an implicit barrier at the end of the sections region/block

# Example: Sections

```
main() {  
    int a[1000], b[1000];  
  
    #pragma omp parallel  
    {  
        #pragma omp sections  
        {  
            #pragma omp section  
            for (int i=0; i<1000; i++)  
                a[i] = 100;  
            #pragma omp section  
            for (int i=0; i<1000; i++)  
                b[i] = 200;  
        }  
    }  
}
```

# Example: Sections (shortcut)

```
main() {  
    int a[1000], b[1000];  
  
    #pragma omp parallel sections  
    {  
  
        #pragma omp section  
        for (int i=0; i<1000; i++)  
            a[i] = 100;  
  
        #pragma omp section  
        for (int i=0; i<1000; i++)  
            b[i] = 200;  
  
    }  
}
```

# Work Sharing in a Parallel Region

Need easy way to split tasks among threads

Most important construct: loops

```
main () {  
    int a[100];  
    #pragma omp parallel  
    {  
        #pragma omp for  
        for (int i= 1; i<n;i++)  
            a[i] = i;  
    }  
}
```

```
main () {  
    int a[100];  
  
    #pragma omp parallel for  
        for (int i= 1; i<n;i++)  
            a[i] = i;  
}
```

Creates a parallel region

Splits iterations of `for` loop among threads

# Parallel Loop

```
#pragma omp for [parameters]  
  for ...
```

- The iterations of the do-loop are distributed to the threads
- There is no synchronization at the beginning
- All threads of the team synchronize at an implicit barrier
  - Unless the parameter **nowait** is specified
- Note: the expressions in the for-statement are very restricted

Iterations must be independent

- No data dependencies
- Can be executed in any order
- Programmer responsibility



# How do Loops get Split up?

Iterations must be distributed to threads

## Loop Schedule

- Defines how iterations are split up
  - Leads to the creation of “Chunks”
- Defines how chunks are mapped to threads

OpenMP offers several options

- Specified using the `schedule` parameter

Example:

```
#pragma omp for schedule(static)
```

# Available Loop Schedules

## static

- Fix sized chunks (default size is about  $n/t$ )
- Distributed in a round-robin fashion

## dynamic

- Fix sized chunks (default size is 1)
- Distributed one by one at runtime as chunks finish

## guided

- Start with large chunks, then exponentially decreasing size
- Distributed one by one at runtime as chunks finish

## runtime

- Controlled at runtime using control variable

## auto

- Compiler/Runtime can choose

# Examples: Scheduling

```
#define S 25
int main(int argc, char** argv)
{ int a[S],b[S],c[S];

#pragma omp parallel
{
    #pragma omp for schedule(static)
    for (int i=0; i<S;i++)
        a[i] = omp_get_thread_num();

    #pragma omp for schedule(dynamic, 4)
    for (int i=0; i<S;i++)
        b[i] = omp_get_thread_num();

    #pragma omp for schedule(guided)
    for (int i=0; i<S;i++)
        c[i] = omp_get_thread_num();
}

for (int i=0; i<S;i++)
    printf("Iter %4d: %4d %4d %4d\n",i,a[i],b[i],c[i]);
}
```


Iter	0:	0	3	2
Iter	1:	0	3	2
Iter	2:	0	3	2
Iter	3:	0	3	1
Iter	4:	0	1	1
Iter	5:	0	1	0
Iter	6:	0	1	0
Iter	7:	1	1	1
Iter	8:	1	0	1
Iter	9:	1	0	0
Iter	10:	1	0	0
Iter	11:	1	0	2
Iter	12:	1	1	1
Iter	13:	2	1	0
Iter	14:	2	1	2
Iter	15:	2	1	0
Iter	16:	2	3	2
Iter	17:	2	3	1
Iter	18:	2	3	0
Iter	19:	3	3	2
Iter	20:	3	0	0
Iter	21:	3	0	2
Iter	22:	3	0	1
Iter	23:	3	0	0
Iter	24:	3	1	2

# Lecture IN-2147

## Parallel Programming

Martin Schulz  
Technical University of Munich  
Department of Computer Engineering

Segment 13  
Data Sharing in OpenMP



# Shared vs. Private Data

Important decision: visibility of data / variables

## Shared data

- Accessible by all threads
- One copy for all threads
- Example: a reference `a[5]` to a shared array accesses

## Private data

- Accessible only by a single thread
- Each thread has its own copy
- Example: iterator variables

The default for global variables: shared

Variables declared within the “dynamic extent” of parallel regions: local

- Includes routines called from parallel regions

# Private clause for parallel loop

```
main () {  
    int a[100], t;  
  
    #pragma omp parallel  
    {  
        #pragma omp for private(t)  
        for (int i= 1; i<n;i++){  
            t=f(i);  
            a[i]=t;  
        }  
    }  
}
```

Global variable: a, t

Local variable: i

Variable t by default shared, but now “privatized”

# Example / Thread Identities

```
main () {
    int iam, nthreads;

    #pragma omp parallel private(iam,nthreads)
    {
        iam = omp_get_thread_num();
        nthreads = omp_get_num_threads();
        printf("ThradID %d, out of %d threads\n", iam, nthreads);

        if (iam == 0)
            printf("Here is the Master Thread.\n");
        else
            printf("Here is another thread.\n");
    }
}
```

# Example / Thread Identities – Better Option

```
main () {  
  
#pragma omp parallel  
{  
    int iam = omp_get_thread_num();  
    int nthreads = omp_get_num_threads();  
    printf("ThreadID %d, out of %d threads\n", iam, nthreads);  
  
    if (iam == 0)  
        printf("Here is the Master Thread.\n");  
    else  
        printf("Here is another thread.\n");  
}  
}
```



# Private Data Options

```
int i=3;
#pragma omp parallel for
    for (int j=0; j<4; j++)
        { i=i+1;
          printf("-> i=%d\n", i); }
printf("Final Value of I=%d\n", i);
```

Before	Loop	After
i = 3		i = 7
	i = 5	
	i = 4	
	i = 7	
	i = 6	

i = 3		i = 4
	i = 4	
	i = 5	
	i = 5	
	i = 6	

```
int i=3;
#pragma omp parallel for private(i)
    for (int j=0; j<4; j++)
        { i=i+1;
          printf("-> i=%d\n", i); }
printf("Final Value of I=%d\n", i);
```

i = 3	i = 3	i = 3
	i = ?	
	i = ?	
	i = ?	
	i = ?	

# First/Last Private Data Options

```
int i=3;
#pragma omp parallel for firstprivate(i)
    for (int j=0; j<4; j++)
    { i=i+1;
      printf("-> i=%d\n", i); }
printf("Final Value of I=%d\n", i);
```

Before	Loop	After
i = 3	i = 3	i = 3
	i = 4	
	i = 4	
	i = 4	
	i = 4	

```
int i=3;
#pragma omp parallel for lastprivate(i)
    for (int j=0; j<4; j++)
    { i=i+1;
      printf("-> i=%d\n", i); }
printf("Final Value of I=%d\n", i);
```

i = 3	i = 3	i = ?
	i = ?	
	i = ?	
	i = ?	
	i = ?	

```
int i=3;
#pragma omp parallel for firstprivate(i) \
                        lastprivate(i)
    for (int j=0; j<4; j++)
    { i=i+1;
      printf("-> i=%d\n", i); }
printf("Final Value of I=%d\n", i);
```

i = 3	i = 3	i = 4
	i = 4	
	i = 4	
	i = 4	
	i = 4	

# Summary: Sharing Attributes of Variables

## `private(var-list)`

- Variables in var-list are private

## `shared(var-list)`

- Variables in var-list are shared

## `default(private | shared | none)`

- Sets the default for all variables in this region

## `firstprivate(var-list)`

- Variables are private
- Initialized with the value of the shared copy before the region.

## `lastprivate(var-list)`

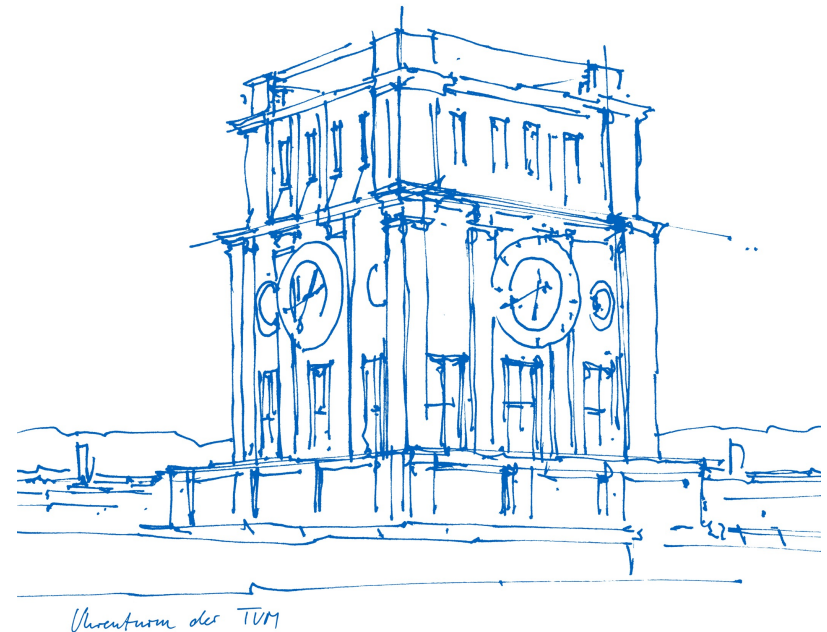
- Variables are private
- Value of the thread executing the last iteration of a parallel loop in sequential order is copied to the variable outside of the region.

# Lecture IN-2147

## Parallel Programming

Martin Schulz  
Technical University of Munich  
Department of Computer Engineering

Segment 14  
Synchronization



# Data Synchronization

As with Pthreads, we need options to synchronize data accesses

- Communication is through shared variables
- Synchronization has to be separate

OpenMP offers a wide range of pragma based options

- Barriers
- Masked regions
- Single regions
- Critical sections
- Atomic statements
- Ordering construct

Additional runtime routines for locking

# Barrier

## Key synchronization construct

- Synchronizes all the threads in a team
- Each thread waits until all other threads in the team have reached the barrier

## Each parallel region has an implicit barrier at the end

- Synchronizes the end of the region
- Can be switched off by adding **nowait**

## Additional barriers can be added when needed

**#pragma omp barrier**

## Warning

- Can cause load imbalance
- Use only when really needed

# Master Region

```
#pragma omp master  
block
```

A master region enforces that only the master executes the code block

- Other threads skip the region
- No synchronization at the beginning of region

Possible uses

- Printing to screen
- Keyboard entries
- File I/O



Warning:  
Deprecated

# Masked Region

```
#pragma omp masked  
    block
```

The **masked** construct (without arguments) declares a region that only the **primary** thread executes the code block

- Other threads skip the region
- No synchronization at the beginning of region

Possible uses

- Printing to screen
- Keyboard entries
- File I/O



# Masked Region

```
#pragma omp masked [filter(integer-expression)]  
    block
```

The **masked** construct (with arguments) declares a region that only the threads execute that are specified in the integer expression (relative to the current team)

- Other threads skip the region
- No synchronization at the beginning of region

## Possible Uses

- Designate a different thread than the primary one for I/O
- Possibly combined with a “nowait” → can create overlap

# Single Region

```
#pragma omp single [parameter]  
    block
```

A single region enforces that only a (arbitrary) single thread executes the code block

- Other threads skip the region
- Implicit barrier synchronization at the end of region (unless **nowait** is specified)

Possible uses

- Initialization of data structures

# Critical Section

```
#pragma omp critical [(Name)]  
    block
```

## Mutual exclusion

- A critical section is a block of code
- Can only be executed by only one thread at a time.
- Compare to Pthreads locks

## Critical section name identifies the specific critical section

- A thread waits at the beginning of a critical section until its available
- All unnamed critical directives map to the same name

## Keep in mind

- Critical section names are global entities of the program
- If a name conflicts with any other entity, program behavior is unspecified
- Avoid long critical sections for performance reasons

# Atomic Statements

```
#pragma ATOMIC  
    expression-stmt
```

The ATOMIC directive ensures that a specific memory location is updated atomically

Has to have the following form:

- `x binop= expr`
- `x++` or `++x`
- `x--` or `--x`
- where `x` is an lvalue expression with scalar type
- and `expr` does not reference the object designated by `x`

Equivalent to using critical section to protect the update

Useful for simple/fast updates to shared data structures

- Avoids locking
- Often implemented directly by native instructions

# Simple Runtime Locks

In addition to pragma based options, OpenMP also offers runtime locks

- Same concept as Pthread mutex
- Locks can be held by only one thread at a time.
- A lock is represented by a lock variable of type `omp_lock_t`.

## Operations

**`omp_init_lock(&lockvar)`**

initialize a lock

**`omp_destroy_lock(&lockvar)`**

destroy a lock

**`omp_set_lock(&lockvar)`**

set lock

**`omp_unset_lock(&lockvar)`**

free lock

**`logicalvar = omp_test_lock(&lockvar)`** check lock and possibly set lock

*returns true if lock was set by the executing thread.*

# Example: Simple Lock

```

#include <omp.h>
int id;
omp_lock_t lock;

omp_init_lock(lock);
#pragma omp parallel shared(lock) private(id)
{
    id = omp_get_thread_num();
    omp_set_lock(&lock); //Only a single thread writes
    printf("My Thread num is: %d", id);
    omp_unset_lock(&lock);
}

while (!omp_test_lock(&lock))
{
    other_work(id); //Lock not obtained
    real_work(id); //Lock obtained
    omp_unset_lock(&lock); //Lock freed
}

omp_destroy_lock(&lock);

```

# Nestable Locks

Similar to simple locks

But, nestable locks can be set multiple times by a single thread.

- Each set operation increments a lock counter
- Each unset operation decrements the lock counter

If the lock counter is 0 after an unset operation, lock can be set by another thread

Separate routines for nestable locks

*Look them up ☺*

# Ordered Construct

```
#pragma omp for ordered
for (...)
{ ...
    #pragma omp ordered
    { ... }
    ...
}
```

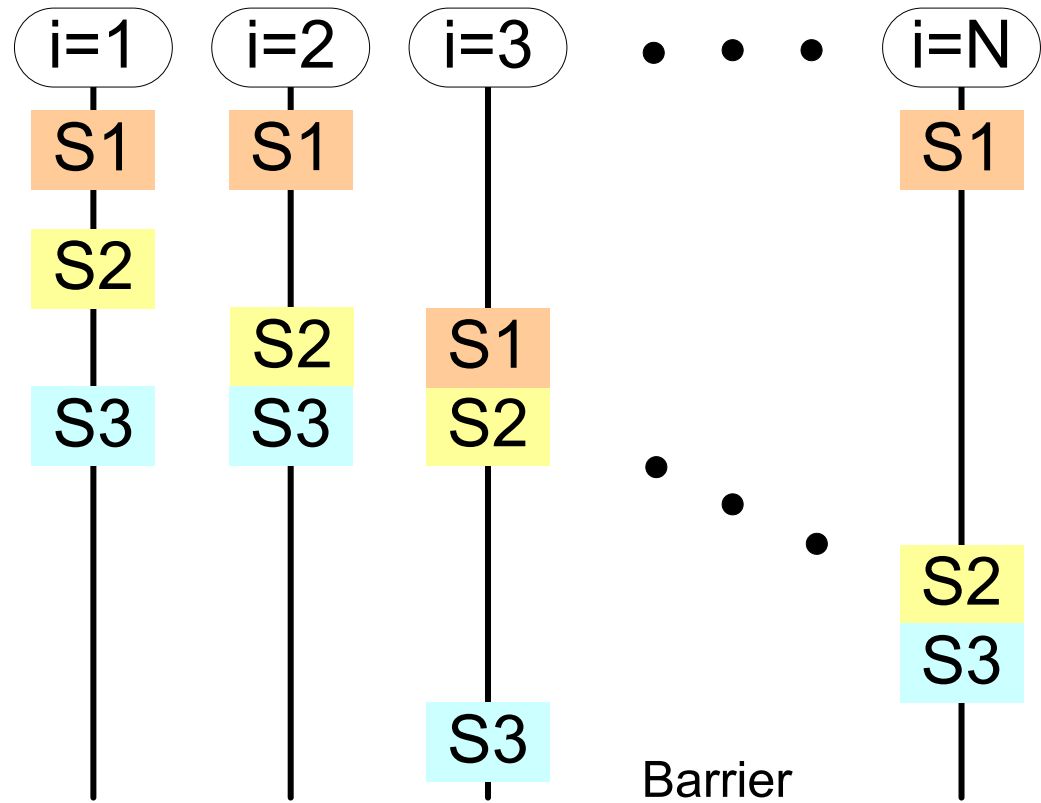
Construct must be within the dynamic extent of an **omp for** construct with an ordered clause.

Ordered constructs are executed strictly in the order in which they would be executed in a sequential execution of the loop.



# Example: ordered clause

```
#pragma omp for ordered
for (...)
{ S1
  #pragma omp ordered
  { S2 }
  S3
}
```



# Summary: OpenMP 101

OpenMP was created to standardize the programming of shared memory systems

- First standard in 1997, currently at OpenMP 5.0
- Goals were easy of use, simplicity and portability

## Key concepts

- Parallel regions
- Worksharing through parallel for loops
- Additional clauses to control distribution, synchronization, ...
- Options to control data visibility/sharing

## Programmer responsibility

- Ensure no loop dependencies exist
- Ensure the right variables are private or shared
- Ensure the necessary synchronization is added